

# A2\_part1\_LSTM\_on\_EHR\_structured

October 30, 2021

## 1 Assignment 2 - clinical prediction with LSTMs on structured MIMIC-III data

The Intensive Care Unit (ICU) treats, on estimate, 55,000 patients per day. ICU patients have an average length of stay of 3.8 days with a mortality rate of 10-29% [source](#). Furthermore, monitoring patients in ICU rooms requires keeping track of tremendous amounts of real-time information, and much of it is logged and stored in electronic health record (EHR) systems. Information overload can be seen as a huge barrier to safe and efficient healthcare delivery. As such, there has been much work on exploring computer assisted diagnostic (CAD) systems to predict clinical outcomes from these data sources.

In the first part of this assignment, you'll be working with structured data measurements (e.g. heart rate, glucose levels, central venous pressure) to make predictions of - sepsis - myocardial infarction (MI) - vancomycin antibiotic administration

over two week patient ICU courses. We'll be running a simplified version of the models in [An attention based deep learning model of clinical events in the intensive care unit](#). (In parts 2 & 3 of this assignment you'll use unstructured clinical text from discharge summaries.)

### Q1.1 clinical applications of prediction models

Many researchers work on the problem of predicting sepsis. Briefly explain how a sepsis prediction model can improve clinical outcomes for patients.

*Written answer:* Accurate early diagnosis of sepsis can reduce the risk of adverse patient outcomes from severe sepsis and septic shock. By identifying systemic inflammation as a sign of infection and detecting possible organ dysfunction, we can begin to recognize early signs of severe sepsis. With advanced warning of impending sepsis onset, clinicians can implement early intervention treatment to drastically improve clinical outcomes for patients.

### 1.1 MIMIC-III Data Preprocessing & Visualization

In the below cell, define the path to `ROOT`. This is where all assignment 2 data will be placed. E.g. you could put it in the same directory as the notebook.

Go to <https://physionet.org/sign-dua/mimiciii/1.4/> and accept the MIMIC-III data use agreement.

Navigate to `ROOT` in your terminal, and then download the MIMIC-III dataset by executing the following commands (replace username with your physionet username). This will create a directory in `ROOT` called `mimic_database/`.

```
wget -r -N -c -np --user <username> --ask-password https://physionet.org/files/mimiciii/1.4/
(2-5 minutes)
```

```
mkdir -p mimic_database && mv physionet.org/files/mimiciii/1.4/*.csv.gz
mimic_database/ && rm -rf physionet.org/ && cd mimic_database && gunzip *.gz &&
echo 'Success!' || echo 'Failure' (5-10 mins)
```

If everything worked, the final output should read **Success!**, and `mimic_database/` should contain some csv files.

```
[1]: import pickle
import math
import re
import csv
import concurrent.futures
import os
from functools import reduce
import pathlib
import pickle

from operator import add
import pandas as pd
import numpy as np

import gc
from time import time
import math
import pickle
import pathlib
import matplotlib.pyplot as plt

import tensorflow as tf
tf.keras.backend.set_floatx('float64')

# packages from current directory
import parser_utils
import data_utils

# config
tf.keras.backend.set_floatx('float32')

ROOT = "/home/marchuo/assign2" # Put your root path here
```

(~40 mins) Run the next cell after setting `DO_PARSING` and `DO_BUILD_DATASETS` to `True`. It's slow, but only needs to be run once. It will create files in `ROOT/mapped_events/`, and in `ROOT/saved_data`, and we'll explain what it's doing a bit later.

Once it runs successfully, set `DO_PARSING` and `DO_BUILD_DATASETS` to `False`.

```
[2]: DO_PARSING=False
DO_BUILD_DATASETS=False

if DO_PARSING: parser_utils.do_all_parsing(ROOT, verbose=1)
if DO_BUILD_DATASETS: data_utils.build_seq_datasets(ROOT)
```

## Q1.2 mimic database exploration I

Let's get a better understanding of the MIMIC-III database. [Here is the documentation](#). The 'Data Description' section is especially useful. When you ran the data download commands at the start of this notebook (the command starting with `wget`), a directory was created in `ROOT` called `mimic_database/`. This contains the MIMIC csv files.

First load the `PATIENTS.csv` file and display the results to screen (hint: use Pandas to load the csv's to a DataFrame; hint 2: the function `display(df)` prints the dataframes nicely).

```
[3]: # YOUR CODE HERE #
mimic_database = "/home/marchuo/assign2/mimic_database/PATIENTS.csv"
patients_df = pd.read_csv(mimic_database)
display(patients_df)
# END CODE #
```

	ROW_ID	SUBJECT_ID	GENDER	DOB	DOD \
0	234	249	F	2075-03-13 00:00:00	NaN
1	235	250	F	2164-12-27 00:00:00	2188-11-22 00:00:00
2	236	251	M	2090-03-15 00:00:00	NaN
3	237	252	M	2078-03-06 00:00:00	NaN
4	238	253	F	2089-11-26 00:00:00	NaN
...	...	...	...	...	...
46515	31840	44089	M	2026-05-25 00:00:00	NaN
46516	31841	44115	F	2124-07-27 00:00:00	NaN
46517	31842	44123	F	2049-11-26 00:00:00	2135-01-12 00:00:00
46518	31843	44126	F	2076-07-25 00:00:00	NaN
46519	31844	44128	M	2098-07-25 00:00:00	NaN

	DOD_HOSP	DOD_SSN	EXPIRE_FLAG
0	NaN	NaN	0
1	2188-11-22 00:00:00	NaN	1
2	NaN	NaN	0
3	NaN	NaN	0
4	NaN	NaN	0
...	...	...	...
46515	NaN	NaN	0
46516	NaN	NaN	0
46517	2135-01-12 00:00:00	NaN	1
46518	NaN	NaN	0
46519	NaN	NaN	0

[46520 rows x 8 columns]

Notice that the date of birth (DOB) and date of death (DOD) are in the future. Briefly explain why (hint: see Methods section of MIMIC documentation).

*Written answer:* Before data could be incorporated into the database, it has to be deidentified. In particular, the dates were deidentified by shifting the dates into the future by some random interval that is consistent with each patient to preserve intervals. This results in stays that occurred in the future between 2100 and 2200.

### Q1.3 mimic database exploration II

In the next code cell, use the dataframe from `PATIENTS.csv` to print the following summary measurements about the dataset: - The number of total patients. - The counts of male and female patients. - The count of patients with a death on record.

(Hint: the `groupby()` function may be useful).

```
[4]: # YOUR CODE HERE #
num_patients = len(patients_df)
num_gender = patients_df.groupby(["GENDER"]).size()
num_deaths = patients_df[patients_df["EXPIRE_FLAG"] == 1]
print("Number of Patients: " + str(num_patients))
print("Number of Patient Deaths: " + str(num_deaths.shape[0]))
print(num_gender)
# END CODE #
```

```
Number of Patients: 46520
Number of Patient Deaths: 15759
GENDER
F      20399
M      26121
dtype: int64
```

### Q1.4 mimic database exploration III

Let's now look at `CHARTEVENTS.csv`, which has one row for each recorded chart measurement (e.g. features like heart rate, glucose levels, central venous pressure). This file is 33GB so don't try to load the whole thing into memory.

Read the first 100 rows of the file `CHARTEVENTS.csv` and display the DataFrame in the notebook.

```
[5]: first_nrows = 100
# YOUR CODE HERE #
events_path = "/home/marchuo/assign2/mimic_database/CHARTEVENTS.csv"
events_df = pd.read_csv(events_path, nrows=first_nrows)
display(events_df)
# END #
```

	ROW_ID	SUBJECT_ID	HADM_ID	ICUSTAY_ID	ITEMID	CHARTTIME	\
0	788	36	165660	241249	223834	2134-05-12 12:00:00	
1	789	36	165660	241249	223835	2134-05-12 12:00:00	
2	790	36	165660	241249	224328	2134-05-12 12:00:00	
3	791	36	165660	241249	224329	2134-05-12 12:00:00	

4	792	36	165660	241249	224330	2134-05-12 12:00:00		
..	...	...	...	...	...	...		
95	348	34	144319	290505	226873	2191-02-23 07:31:00		
96	349	34	144319	290505	220210	2191-02-23 07:33:00		
97	350	34	144319	290505	220045	2191-02-23 07:34:00		
98	351	34	144319	290505	220179	2191-02-23 07:34:00		
99	352	34	144319	290505	220180	2191-02-23 07:34:00		

	STORETIME	CGID	VALUE	VALUENUM	VALUEUOM	WARNING	ERROR	\
0	2134-05-12 13:56:00	17525	15.00	15.00	L/min	0	0	
1	2134-05-12 13:56:00	17525	100.00	100.00	NaN	0	0	
2	2134-05-12 12:18:00	20823	0.37	0.37	NaN	0	0	
3	2134-05-12 12:19:00	20823	6.00	6.00	min	0	0	
4	2134-05-12 12:19:00	20823	2.50	2.50	NaN	0	0	
..	...	...	...	...	...	...	...	
95	2191-02-23 07:35:00	16924	1.00	1.00	NaN	0	0	
96	2191-02-23 07:45:00	17741	26.00	26.00	insp/min	0	0	
97	2191-02-23 10:53:00	17741	44.00	44.00	bpm	0	0	
98	2191-02-23 07:45:00	17741	135.00	135.00	mmHg	0	0	
99	2191-02-23 07:45:00	17741	61.00	61.00	mmHg	0	0	

	RESULTSTATUS	STOPPED
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
..	...	...
95	NaN	NaN
96	NaN	NaN
97	NaN	NaN
98	NaN	NaN
99	NaN	NaN

[100 rows x 15 columns]

Each row is a single measurement, but it does not say what is being measured. Explain how you could find this out. (Hint: see the ‘Data Description’ section of the documentation.)

*Written answer:* Tables are linked by identifiers which have ‘ID’ as a suffix. For example, SUBJECT\_ID refers to a patient, ICUSTAY\_ID refers to a specific ICU visit, and HADM\_ID refers to a specific hospital visit. Charted events are stored as a series of events tables, such as OUTPUTEVENTS and LABEVENTS. By joining dictionary tables that store identifiers prefixed with “D\_” and events tables on ITEMID, we can find out what is being measured. Essentially we are cross-referencing codes stored in dictionary tables against their respective definitions in the events tables.

So far we’ve looked at the original MIMIC-III database, but it’s not in a format suitable for a sequence model prediction (like LSTMs or transformers). You ran two lines of code at the start of

the assignment to get it in the right format.

The first function was `parser_utils.do_all_parsing`, which created a set of files in `ROOT/mimic_database/mapped_events`, which are closer to what we need. One output was the file `CHARTEVENTS_reduced_24_hour_blocks_plus_admissions_plus_patients_plus_scripts_plus_icds_plus_n`. This file: - Is similar to `CHARTEVENTS`, except that each measurement is **assigned to a single 24hr block** (like 2117-09-11), rather than a specific timestamp (like 2117-09-11 16:04:00). You can think of this as discretizing the dataset. - Each row also has extra data about the patient: admission times, scripts, and known patient diseases (ICD's).

Next you ran `data_utils.load_seq_dataset` which does the following: - Given a prediction target (one of `MI`, `SEPSIS`, or `VANCOMYCIN`), generate numpy arrays for the train, test, and validation set. The data is shuffled before splitting. - Put X-data into the shape `(n_hostpital_stays, n_timesteps, n_features)`. So `X[i,j,k]` gives the `k`th feature, for the `j`th day of the `i`th hospital stay. - Puts the labels, y-data, into shape `(n_hospital_stays,n_timesteps,1)`. All 3 prediction problems are binary, to these values 0 or 1. - Returns a list of strings called `features` containing the names of each feature in the X-data. So `features[k]` is the name of the features in `X[:, :, k]` - Zero-padding. Since not all patients will have a valid measurement for every feature at each timestep, we fill the remainder with zeros. Later we will tell the model to ignore this data in training. - Z-score normalization.

Now we can load that data using the following function call:

```
[6]: target='SEPSIS'      # 'SEPSIS' or 'MI' or 'VANCOMYCIN'
train_x, val_x, train_y, val_y, no_feature_cols, test_x, test_y, \
    →x_boolmat_test, y_boolmat_test, x_boolmat_val, y_boolmat_val, features \
    = data_utils.load_seq_dataset(ROOT, target)

print("train shapes ", train_x.shape, train_y.shape)
print("val shapes   ", val_x.shape, val_y.shape)
print("test shapes  ", test_x.shape, test_y.shape)
print("# features   ", len(features))
```

```
train shapes  (2984, 15, 226) (2984, 15, 1)
val shapes    (5178, 15, 226) (5178, 15, 1)
test shapes   (10355, 15, 226) (10355, 15, 1)
# features    226
```



### Q1.5 sequence model data sets



Let's look at some samples from `train_x`. In the below code cell, get the feature called 'HDL', and generate 20 plots showing how this variable changes over all the timesteps for the first 20 hospital stays (1 time series plot per hospital stay).

```
[7]: # YOUR CODE HERE #
num_stays = 20
hdl_i = features.index("HDL")
for i in range(num_stays):
    i_data = train_x[i, :, hdl_i]
    plt.plot(i_data)
```



```
plt.show()
```



```
# END CODE #
```



































By looking at these graphs and by carefully reading the description of `return_data` above, answer the following questions: - Why are the last values in the series usually 0, and why do some reach 0 earlier than others? - “HDL” is cholesterol, so how is it possible that some values are actually negative? - In many cases, the measured HDL cholesterol is exactly the same on consecutive days, even though we would expect a measurement to fluctuate at least a bit between days. Why is this



the case?

*Written answers* - . The last values in the series are usually 0, which means that no data has been collected for that timestep. Some tables reach 0 earlier than others, as the patient may have been released from the hospital earlier than others, which cuts out the data available beyond that timestep. In addition, generally the data may not have been entered manually by the attending clinician. - . The HDL is negative as the value is z-score normalized so every value above the mean is positive and below the mean is negative. - . The measured HDL cholesterol values may be exactly the same on consecutive days as a result of clinician human error and inputting the same cholesterol values across the following hours. We need a value for every timestep, so we re-use the same result from the previous hour and as the current hour without measuring it.

## Q1.6 prediction problem

In the first cell of this notebook we explained the prediction problem at a high level. Explain the prediction problem again, but be specific in explaining the data inputs and prediction outputs. Using the terminology from lecture, is this a “many-to-many” problem, or a “many-to-one” problem?

*Written Answer:* By working with structured data measurements such as heart rate, glucose, and central venous pressures as our input values, we’ll be making early predictions of sepsis, myocardial infarction, and vancomycin antibiotic administration. As such, this is a many-to-many problem as we are taking multiple structured data measurements as our inputs and outputting one prediction at each timestep of whether the patient is at risk of sepsis, myocardial infarction, and vancomycin antibiotic administration.

## 1.2 LSTM prediction model

### Q1.7 define an LSTM

We’ll use an LSTM model to predict a patient outcome at each time step. So each data point we pass to the model will be a sequence of length `n_timesteps`, where each timestamp has `k` features. Our prediction output is also a sequence of length `n_timesteps`.

Using Keras `Sequential`, define a model called `model_lstm`. It should have: - 1 LSTM layer with 256 units. Use the default activation for the LSTM layer. Keras has optimized GPU implementations for most layers, but it does not have an optimized implementation for LSTM with non-default activations. - 1 dropout layer with `rate=0.5`. - 1 dense layer that applies the same transformation to each of the prediction outputs. - A suitable activation function for the prediction task.

```
[8]: def build_lstm_model(lstm_hidden_units= 256):  
    """  
    Return a simple Keras model with a single LSTM layer, dropout later,  
    and then dense prediction layer.  
  
    Args:  
    lstm_hidden_units (int): units in the LSTM layer  
  
    Returns:  
    model_lstm (tf.keras.Model) LSTM keras model with output dimension (None,1)  
    """
```

```

model_lstm = None
# YOUR CODE HERE #
model_lstm = tf.keras.Sequential(
    layers=[
        tf.keras.layers.LSTM(lstm_hidden_units, return_sequences=True),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(1, activation='sigmoid')
    ]
)
# END CODE #
return model_lstm

# test code for checking the shape #
lstm_hidden_units = 256
model_lstm = build_lstm_model(lstm_hidden_units)
bs=8
x_batch = train_x[:bs]
print(model_lstm(x_batch).shape) # expect shape (8, 15, 1) # batch size 8, 15,
    ↳ timesteps

```

```

2021-10-30 02:05:01.971847: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:01.982831: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:01.983596: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:01.986266: I tensorflow/core/platform/cpu_feature_guard.cc:142]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
2021-10-30 02:05:01.986719: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:01.987741: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:01.988505: I

```

```

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:02.395285: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:02.396099: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:02.396857: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:05:02.397583: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 10819 MB memory:  -> device:
0, name: Tesla K80, pci bus id: 0000:00:04.0, compute capability: 3.7
2021-10-30 02:05:03.027770: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369]
Loaded cuDNN version 8005

```

(8, 15, 1)

### Q1.8 masking

Read about the [masking layer](#) in Keras. Briefly explain why we need masking for this problem. Your answer should refer back to the time-series plots generated in Q1.5.

*Written answer:* The masking layer is used to mask a sequence by using a mask value to skip timesteps. We need masking for this problem, because as shown in the time-series plots generated in Q1.5, we have missing timesteps in our input, which we need to skip when processing the data as that would affect our final model prediction accuracy and efficiency.

The below function builds the final model. We provide code that calls `build_lstm_model`.

Your code should add a masking layer with `mask_value=0`, and it should be applied at the start of the model.

```

[9]: def build_masked_lstm_model(num_timesteps, num_features, lstm_hidden_units=256):
    """
    Return a simple Keras model with a masking single LSTM layer, dropout_
    ↪ later,
    and then dense prediction layer.

    Args:
    num_timesteps (int): num timesteps per input data object.
    num_features (int): num features per input data object.
    lstm_hidden_units (int): units in the LSTM layer
    """

```

```

Returns:
model_lstm (tf.keras.Model) LSTM keras model with output dimension (None,1)
"""

model_lstm = build_lstm_model(lstm_hidden_units)
for layer in model_lstm.layers:
    layer.supports_masking=True

model = None
# YOUR CODE HERE #
model = tf.keras.Sequential()
model.add(tf.keras.layers.Masking(mask_value=0.,
                                   input_shape=(num_timesteps, num_features)))

model.add(model_lstm)
# END CODE #
return model

# Code to test the shape is correc #
num_timesteps, num_features = train_x.shape[-2:]
lstm_hidden_units = 256

model = build_masked_lstm_model(num_timesteps, num_features, lstm_hidden_units)
bs=8
x_batch = train_x[:bs]
print(model(x_batch).shape) # expect shape (8, 15, 1) # batch size 8, 15,
→timesteps

```

(8, 15, 1)

### Q1.9 compiling and training

Below we have copied the code for getting the dataset that we ran earlier. You can choose ‘MI’, SEPSIS’ or ‘VANCOMYCIN’ as the target. When submitting this assignment, please choose ‘VANCOMYCIN’.

The code also calls `build_masked_lstm_model` that you just defined. Note that the model parameters depend on the dataset shape, so if we wanted to change the dataset from ‘MI’ to ‘SEPSIS’ then we need to create the model again with different input shapes. We chose to define the model inside a function so that we could re-create the model more easily.

Compile the model using - The Adam optimizer with default parameters. - An appropriate loss function for this task. - Metrics: accuracy and tensorflow’s [AUC](#)

```

[10]: target='VANCOMYCIN' # 'SEPSIS' or 'MI' or 'VANCOMYCIN'
train_x, val_x, train_y, val_y, no_feature_cols, test_x, test_y,
→x_boolmat_test, y_boolmat_test, x_boolmat_val, y_boolmat_val, features \
    = data_utils.load_seq_dataset(ROOT, target)
num_timesteps, num_features = train_x.shape[-2:]
model = build_masked_lstm_model(num_timesteps, num_features, lstm_hidden_units)

```

```
# YOUR CODE HERE #
model.compile(loss=tf.keras.losses.BinaryCrossentropy(), optimizer=tf.keras.
    ↳optimizers.Adam(learning_rate=0.001), metrics=['accuracy', tf.keras.metrics.
    ↳AUC()])
# YOUR CODE HERE #
```

Finally fit the model. It should train very quickly. For ‘VANCOMYCIN’, validation accuracy should be around 0.85. For ‘MI’ it should be above 0.95.

```
[11]: epochs=10
# YOUR CODE HERE #
hist = model.fit(x=train_x, y=train_y, epochs=epochs, validation_data=(val_x,
    ↳val_y))
# YOUR CODE HERE #
```

```
2021-10-30 02:05:07.346604: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR
Optimization Passes are enabled (registered 2)
```

```
Epoch 1/10
914/914 [=====] - 22s 19ms/step - loss: 0.2413 -
accuracy: 0.8110 - auc: 0.8698 - val_loss: 0.1933 - val_accuracy: 0.8413 -
val_auc: 0.9005
Epoch 2/10
914/914 [=====] - 16s 17ms/step - loss: 0.2143 -
accuracy: 0.8337 - auc: 0.8995 - val_loss: 0.1860 - val_accuracy: 0.8484 -
val_auc: 0.9088
Epoch 3/10
914/914 [=====] - 16s 17ms/step - loss: 0.2055 -
accuracy: 0.8413 - auc: 0.9081 - val_loss: 0.1802 - val_accuracy: 0.8541 -
val_auc: 0.9149
Epoch 4/10
914/914 [=====] - 16s 17ms/step - loss: 0.1986 -
accuracy: 0.8463 - auc: 0.9146 - val_loss: 0.1828 - val_accuracy: 0.8528 -
val_auc: 0.9119
Epoch 5/10
914/914 [=====] - 16s 17ms/step - loss: 0.1924 -
accuracy: 0.8521 - auc: 0.9202 - val_loss: 0.1797 - val_accuracy: 0.8544 -
val_auc: 0.9161
Epoch 6/10
914/914 [=====] - 16s 17ms/step - loss: 0.1861 -
accuracy: 0.8576 - auc: 0.9257 - val_loss: 0.1784 - val_accuracy: 0.8552 -
val_auc: 0.9172
Epoch 7/10
914/914 [=====] - 16s 17ms/step - loss: 0.1811 -
accuracy: 0.8612 - auc: 0.9298 - val_loss: 0.1801 - val_accuracy: 0.8542 -
val_auc: 0.9165
Epoch 8/10
```

```

914/914 [=====] - 16s 17ms/step - loss: 0.1745 -
accuracy: 0.8665 - auc: 0.9350 - val_loss: 0.1820 - val_accuracy: 0.8526 -
val_auc: 0.9148
Epoch 9/10
914/914 [=====] - 16s 17ms/step - loss: 0.1681 -
accuracy: 0.8735 - auc: 0.9401 - val_loss: 0.1850 - val_accuracy: 0.8531 -
val_auc: 0.9134
Epoch 10/10
914/914 [=====] - 16s 17ms/step - loss: 0.1588 -
accuracy: 0.8817 - auc: 0.9466 - val_loss: 0.1931 - val_accuracy: 0.8503 -
val_auc: 0.9108

```

### 1.2.1 Evaluation

#### Q1.10 prediction and masking

Use `model.predict()` on the test dataset and save predictions to the variable `test_y_pred`.

```

[12]: test_y_pred = None
      # YOUR CODE HERE #
      test_y_pred = model.predict(test_x)
      # END CODE #

      print(test_y_pred.shape) # expect (n_datapoints, 15, 1)

```

```
(10355, 15, 1)
```

Normally when we compare a set of predictions, we'd take 2 vectors: `y_true` which is a flat vector of 0s and 1s, and `y_pred` which is a vector of the same shape with probabilities in the range [0,1]. Our case is different because: - Our prediction output has an extra axis (`None,timesteps,features`) instead of (`None,features`). - Some of the predictions should be masked, and therefore removed from the final prediction dataset.

The earlier function `data_utils.load_seq_dataset` returned a mask vector `y_boolmat_test` with the same shape as `test_y`. If `y_boolmat_test[i]==True` then this label should be masked (removed from the evaluation dataset).

In the next cell use `test_y_pred` and `y_boolmat_test` to create the vectors `y_pred_masked` and `y_true_masked` by removing the masked predictions, and flattening the output.

```

[13]: y_pred_masked = None
      y_true_masked = None

      # YOUR CODE HERE #
      y_pred_masked = []
      y_true_masked = []
      for i in range(y_boolmat_test.shape[0]):
          for n in range(y_boolmat_test.shape[1]):
              if y_boolmat_test[i, n, 0] == False:
                  y_pred_masked.append(test_y_pred[i, n, 0])

```

```

        y_true_masked.append(test_y[i, n, 0])
y_pred_masked = np.array(y_pred_masked)
y_true_masked = np.array(y_true_masked)
# END CODE #
print(y_pred_masked.shape, y_true_masked.shape) # expect shape (n_predictions,)
↳ and the shape should be the same

```

(87373,) (87373,)

### Q1.11 ROC+AUC

Using `y_pred_masked` and `y_true_masked`: - Plot a ROC curve. - Print the AUC.


You can use the functions in `sklearn.metrics` for both.

```

[14]: # YOUR CODE HERE #
from sklearn.metrics import roc_curve, roc_auc_score
fpr, tpr, thresholds = roc_curve(y_true_masked, y_pred_masked)
plt.plot(fpr, tpr)
plt.plot([0, 1], [0, 1], 'k--')
auc = roc_auc_score(y_true_masked, y_pred_masked)
print(f"AUC score={auc}")
# END CODE #

```

AUC score=0.9130357116229655



### Q1.12 confusion


Finally, generate a confusion matrix. To do this you will need to convert prediction probabilities

in `y_pred_masked` to binary predictions. You can choose a threshold of 0.5. Again, you can use functions from `sklearn.metrics`.

You can use a plotting library to display the confusion matrix, but you can also just print the array directly. If you do just print the array, then also print a message explaining the each axis.

```
[15]: # YOUR CODE HERE #
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.preprocessing import binarize
threshold = 0.5
for i in range(y_pred_masked.shape[0]):
    if y_pred_masked[i] < threshold:
        y_pred_masked[i] = 0
    else:
        y_pred_masked[i] = 1
confusion = confusion_matrix(y_true_masked, y_pred_masked)
disp = ConfusionMatrixDisplay(confusion)
disp.plot()
# END CODE #
```

```
[15]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7f9b63fa84d0>
```



### Q1.13 clinical application tradeoffs

We have focused on modelling sepsis, but now consider vancomycin prediction. Explain what it



means to choose different operating points at different positions in the ROC curve. Specifically tie it back to the use case of vancomycin antibiotic administration. What are the tradeoffs? Pick a two points on the ROC curve and explain what the true positive and false positive rates mean at those points.

*Written answer:* To choose different operating points at different positions in the ROC curve means to choose tradeoffs between specificity and sensitivity. An operating point with high sensitivity corresponds to high negative predictive value, which is ideal for “rule-out” tests. On the other hand, an operating point with high specificity corresponds to high positive predictive value, which is ideal for “rule-in” tests. For high specificity, it can be used to test negativity in health, where we want to know the proportion of population without the disease and give negative test results. High specificity should be used to make decisions about high-risk actions. For our specific use case of vancomycin antibiotic administration, it is important to select an effective operating point as a guideline for dosing, as it’s necessary to balance nephrotoxicity with vancomycin’s antibiotic activity. In addition, it’s important to not over-dose antibiotics to reduce incidence of antibiotic resistance. The most optimal operating point to choose is usually when the classifier gives the best trade off between the costs of failing to detect positives against the costs of raising false alarms. The first point on the ROC curve I’ll choose is when both the FPR and TPR are low, when the fpr is equal to 0.05. At this point, the TPR is 0.6, which is a relatively strict threshold to choose. At this point, the true positive rate means how many correct positive results occur among all positive samples available, which is 0.6, while the fpr is 0.05, which defines how many incorrect positive results occur among all negative samples available during the test. The second point on the ROC curve I’ll choose is when FPR is 0.6 and TPR is around 0.95. At this point, the model predicts 0.95 correct positive results among all positive. samples and 0.6 incorrect positive results among all negative samples. At these two extremes, we can see the clear tradeoffs between sensitivity vs. specificity.

[ ]:

# A2\_part2\_clinical\_Word2Vec\_embeddings\_and\_readmission\_prediction

October 30, 2021

## 1 Assignment 2 - part 2 - Clinical Word Embeddings For Prediction

In part 1 you used structured sequence data to make predictions. In part 2 we will ignore that structured data and only use unstructured clinician notes from MIMIC-III. We will use discharge summaries to predict 30-day hospital readmission.

**Importantly** there are two separate distinct steps: 1. Learn good word embeddings. Word embeddings are function that maps words to fixed-length vectors (e.g. 32-dims). We want words that are similar in meaning to similar vector embeddings. 2. Create a deep learning model that takes text input and predicts whether a patient is readmitted. The inputs to the model will be word embeddings from the first step.

We will approach task 1, learning word embeddings, using the popular Word2Vec algorithm (see [original paper](#)). We'll use the skip-gram version of Word2Vec (the other version is 'continuous bag of words')

We will approach task 2 with an LSTM.

### Q2.1

Explain how a model for 30-day readmission prediction could be used by doctors in a clinical setting.

*Written answer: A model for 30-day readmission prediction can be used to reduce readmission risk by assessing doctor notes and predicting whether a patient is readmitted. By predicting whether a patient is readmitted, clinicians can focus on complementing inpatient care with postdischarge interventions and/or enhanced care transition. Early interventions during inpatient hospitalization such as early discharge planning can serve to reduce readmissions. As health care resources are limited, it's essential to predict which patients are at highest risk of readmission and invest readmission-preventative interventions to these specific patients to reduce 30-day readmissions.*

Install the following packages.

```
[1]: !pip install gensim
      !pip install spacy==2.3.7
      !pip install scispacy==0.3.0
      !pip install nltk
      !pip install tdqm
```

```
!pip install https://s3-us-west-2.amazonaws.com/ai2-s2-scispace/releases/v0.2.5/
↪en_core_sci_md-0.2.5.tar.gz
```

Requirement already satisfied: gensim in /opt/conda/lib/python3.7/site-packages (4.1.2)

Requirement already satisfied: smart-open>=1.8.1 in /opt/conda/lib/python3.7/site-packages (from gensim) (5.2.1)

Requirement already satisfied: scipy>=0.18.1 in /opt/conda/lib/python3.7/site-packages (from gensim) (1.7.1)

Requirement already satisfied: numpy>=1.17.0 in /opt/conda/lib/python3.7/site-packages (from gensim) (1.19.5)

Requirement already satisfied: spacy==2.3.7 in /opt/conda/lib/python3.7/site-packages (2.3.7)

Requirement already satisfied: srsly<1.1.0,>=1.0.2 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (1.0.5)

Requirement already satisfied: wasabi<1.1.0,>=0.4.0 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (0.8.2)

Requirement already satisfied: catalogue<1.1.0,>=0.0.7 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (1.0.0)

Requirement already satisfied: thinc<7.5.0,>=7.4.1 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (7.4.5)

Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (4.62.3)

Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (1.0.5)

Requirement already satisfied: cymem<2.1.0,>=2.0.2 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (2.0.5)

Requirement already satisfied: numpy>=1.15.0 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (1.19.5)

Requirement already satisfied: requests<3.0.0,>=2.13.0 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (2.25.1)

Requirement already satisfied: blis<0.8.0,>=0.4.0 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (0.7.4)

Requirement already satisfied: plac<1.2.0,>=0.9.6 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (1.1.3)

Requirement already satisfied: setuptools in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (58.0.4)

Requirement already satisfied: preshed<3.1.0,>=3.0.2 in /opt/conda/lib/python3.7/site-packages (from spacy==2.3.7) (3.0.5)

Requirement already satisfied: importlib-metadata>=0.20 in /opt/conda/lib/python3.7/site-packages (from catalogue<1.1.0,>=0.0.7->spacy==2.3.7) (4.8.1)

Requirement already satisfied: typing-extensions>=3.6.4 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata>=0.20->catalogue<1.1.0,>=0.0.7->spacy==2.3.7) (3.10.0.2)

Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata>=0.20->catalogue<1.1.0,>=0.0.7->spacy==2.3.7)

(3.5.0)

Requirement already satisfied: certifi>=2017.4.17 in  
/opt/conda/lib/python3.7/site-packages (from  
requests<3.0.0,>=2.13.0->spacy==2.3.7) (2021.5.30)  
Requirement already satisfied: urllib3<1.27,>=1.21.1 in  
/opt/conda/lib/python3.7/site-packages (from  
requests<3.0.0,>=2.13.0->spacy==2.3.7) (1.26.6)  
Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-  
packages (from requests<3.0.0,>=2.13.0->spacy==2.3.7) (2.10)  
Requirement already satisfied: chardet<5,>=3.0.2 in  
/opt/conda/lib/python3.7/site-packages (from  
requests<3.0.0,>=2.13.0->spacy==2.3.7) (4.0.0)  
Requirement already satisfied: scispacy==0.3.0 in /opt/conda/lib/python3.7/site-  
packages (0.3.0)  
Requirement already satisfied: nmslib>=1.7.3.6 in /opt/conda/lib/python3.7/site-  
packages (from scispacy==0.3.0) (2.1.1)  
Requirement already satisfied: spacy<3.0.0,>=2.3.0 in  
/opt/conda/lib/python3.7/site-packages (from scispacy==0.3.0) (2.3.7)  
Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages  
(from scispacy==0.3.0) (1.19.5)  
Requirement already satisfied: requests<3.0.0conllu,>=2.0.0 in  
/opt/conda/lib/python3.7/site-packages (from scispacy==0.3.0) (2.25.1)  
Requirement already satisfied: scikit-learn>=0.20.3 in  
/opt/conda/lib/python3.7/site-packages (from scispacy==0.3.0) (0.24.2)  
Requirement already satisfied: joblib in /opt/conda/lib/python3.7/site-packages  
(from scispacy==0.3.0) (1.0.1)  
Requirement already satisfied: pysbd in /opt/conda/lib/python3.7/site-packages  
(from scispacy==0.3.0) (0.3.4)  
Requirement already satisfied: pybind11<2.6.2 in /opt/conda/lib/python3.7/site-  
packages (from nmslib>=1.7.3.6->scispacy==0.3.0) (2.6.1)  
Requirement already satisfied: psutil in /opt/conda/lib/python3.7/site-packages  
(from nmslib>=1.7.3.6->scispacy==0.3.0) (5.8.0)  
Requirement already satisfied: certifi>=2017.4.17 in  
/opt/conda/lib/python3.7/site-packages (from  
requests<3.0.0conllu,>=2.0.0->scispacy==0.3.0) (2021.5.30)  
Requirement already satisfied: chardet<5,>=3.0.2 in  
/opt/conda/lib/python3.7/site-packages (from  
requests<3.0.0conllu,>=2.0.0->scispacy==0.3.0) (4.0.0)  
Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-  
packages (from requests<3.0.0conllu,>=2.0.0->scispacy==0.3.0) (2.10)  
Requirement already satisfied: urllib3<1.27,>=1.21.1 in  
/opt/conda/lib/python3.7/site-packages (from  
requests<3.0.0conllu,>=2.0.0->scispacy==0.3.0) (1.26.6)  
Requirement already satisfied: scipy>=0.19.1 in /opt/conda/lib/python3.7/site-  
packages (from scikit-learn>=0.20.3->scispacy==0.3.0) (1.7.1)  
Requirement already satisfied: threadpoolctl>=2.0.0 in  
/opt/conda/lib/python3.7/site-packages (from scikit-  
learn>=0.20.3->scispacy==0.3.0) (2.2.0)

Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (1.0.5)

Requirement already satisfied: catalogue<1.1.0,>=0.0.7 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (1.0.0)

Requirement already satisfied: srsly<1.1.0,>=1.0.2 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (1.0.5)

Requirement already satisfied: preshed<3.1.0,>=3.0.2 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (3.0.5)

Requirement already satisfied: wasabi<1.1.0,>=0.4.0 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (0.8.2)

Requirement already satisfied: setuptools in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (58.0.4)

Requirement already satisfied: plac<1.2.0,>=0.9.6 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (1.1.3)

Requirement already satisfied: cymem<2.1.0,>=2.0.2 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (2.0.5)

Requirement already satisfied: blis<0.8.0,>=0.4.0 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (0.7.4)

Requirement already satisfied: thinc<7.5.0,>=7.4.1 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (7.4.5)

Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in /opt/conda/lib/python3.7/site-packages (from spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (4.62.3)

Requirement already satisfied: importlib-metadata>=0.20 in /opt/conda/lib/python3.7/site-packages (from catalogue<1.1.0,>=0.0.7->spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (4.8.1)

Requirement already satisfied: typing-extensions>=3.6.4 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata>=0.20->catalogue<1.1.0,>=0.0.7->spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (3.10.0.2)

Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata>=0.20->catalogue<1.1.0,>=0.0.7->spacy<3.0.0,>=2.3.0->scispacy==0.3.0) (3.5.0)

Requirement already satisfied: nltk in /opt/conda/lib/python3.7/site-packages (3.6.5)

Requirement already satisfied: joblib in /opt/conda/lib/python3.7/site-packages (from nltk) (1.0.1)

Requirement already satisfied: tqdm in /opt/conda/lib/python3.7/site-packages

```

(from nltk) (4.62.3)
Requirement already satisfied: regex>=2021.8.3 in /opt/conda/lib/python3.7/site-
packages (from nltk) (2021.8.28)
Requirement already satisfied: click in /opt/conda/lib/python3.7/site-packages
(from nltk) (8.0.1)
Requirement already satisfied: importlib-metadata in
/opt/conda/lib/python3.7/site-packages (from click->nltk) (4.8.1)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-
packages (from importlib-metadata->click->nltk) (3.5.0)
Requirement already satisfied: typing-extensions>=3.6.4 in
/opt/conda/lib/python3.7/site-packages (from importlib-metadata->click->nltk)
(3.10.0.2)
Requirement already satisfied: tqdm in /opt/conda/lib/python3.7/site-packages
(0.0.1)
Requirement already satisfied: tqdm in /opt/conda/lib/python3.7/site-packages
(from tqdm) (4.62.3)
Collecting https://s3-us-
west-2.amazonaws.com/ai2-s2-scispacey/releases/v0.2.5/en_core_sci_md-0.2.5.tar.gz
  Using cached https://s3-us-
west-2.amazonaws.com/ai2-s2-scispacey/releases/v0.2.5/en_core_sci_md-0.2.5.tar.gz
(79.9 MB)
Requirement already satisfied: spacy>=2.3.0 in /opt/conda/lib/python3.7/site-
packages (from en-core-sci-md==0.2.5) (2.3.7)
Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (4.62.3)
Requirement already satisfied: blis<0.8.0,>=0.4.0 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (0.7.4)
Requirement already satisfied: requests<3.0.0,>=2.13.0 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (2.25.1)
Requirement already satisfied: catalogue<1.1.0,>=0.0.7 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (1.0.0)
Requirement already satisfied: thinc<7.5.0,>=7.4.1 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (7.4.5)
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (2.0.5)
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (1.0.5)
Requirement already satisfied: wasabi<1.1.0,>=0.4.0 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (0.8.2)
Requirement already satisfied: srsly<1.1.0,>=1.0.2 in

```

```

/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (1.0.5)
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (3.0.5)
Requirement already satisfied: numpy>=1.15.0 in /opt/conda/lib/python3.7/site-
packages (from spacy>=2.3.0->en-core-sci-md==0.2.5) (1.19.5)
Requirement already satisfied: plac<1.2.0,>=0.9.6 in
/opt/conda/lib/python3.7/site-packages (from spacy>=2.3.0->en-core-sci-
md==0.2.5) (1.1.3)
Requirement already satisfied: setuptools in /opt/conda/lib/python3.7/site-
packages (from spacy>=2.3.0->en-core-sci-md==0.2.5) (58.0.4)
Requirement already satisfied: importlib-metadata>=0.20 in
/opt/conda/lib/python3.7/site-packages (from
catalogue<1.1.0,>=0.0.7->spacy>=2.3.0->en-core-sci-md==0.2.5) (4.8.1)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-
packages (from importlib-
metadata>=0.20->catalogue<1.1.0,>=0.0.7->spacy>=2.3.0->en-core-sci-md==0.2.5)
(3.5.0)
Requirement already satisfied: typing-extensions>=3.6.4 in
/opt/conda/lib/python3.7/site-packages (from importlib-
metadata>=0.20->catalogue<1.1.0,>=0.0.7->spacy>=2.3.0->en-core-sci-md==0.2.5)
(3.10.0.2)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in
/opt/conda/lib/python3.7/site-packages (from
requests<3.0.0,>=2.13.0->spacy>=2.3.0->en-core-sci-md==0.2.5) (1.26.6)
Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-
packages (from requests<3.0.0,>=2.13.0->spacy>=2.3.0->en-core-sci-md==0.2.5)
(2.10)
Requirement already satisfied: certifi>=2017.4.17 in
/opt/conda/lib/python3.7/site-packages (from
requests<3.0.0,>=2.13.0->spacy>=2.3.0->en-core-sci-md==0.2.5) (2021.5.30)
Requirement already satisfied: chardet<5,>=3.0.2 in
/opt/conda/lib/python3.7/site-packages (from
requests<3.0.0,>=2.13.0->spacy>=2.3.0->en-core-sci-md==0.2.5) (4.0.0)

```

Change ROOT to your path.

```

[2]: import os
import tensorflow as tf
import numpy as np
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt

import tqdm
import pandas as pd
import random
import pickle

```

```
import readmission_utils

from tensorflow.keras.preprocessing.text import text_to_word_sequence
from tensorflow.keras.utils import to_categorical

ROOT = "/home/marchuo/assign2" # Put your root path here"
NUM_NS=4 # number of negative samples in Word2Vec model
VOCAB_SIZE=500 # numer of most common words to index in language models
tf.keras.backend.set_floatx('float32')
```

## 1.1 Preprocessing text data and visualization

Execute the code in the next cell, which will take about 20mins the first time you run it. It will save its results to a file in `ROOT/saved_data/texts_to_labels_1000.pkl`.

If the file already exists then calling the function will just load the results. We'll explain what it's doing later.

```
[3]: notes, labels_admission = readmission_utils.get_notes_and_labels(ROOT, 1000)
```

Found file `/home/marchuo/assign2/saved_data/texts_to_labels_1000.pkl`, loading

## Q2.2 admissions database

As in part 1, let's briefly look at the underlying data tables. Our task will be to predict hospital readmission, so we're interested in the file `ADMISSIONS.csv` which is in `ROOT/mimic_database`. Load the table to a dataframe and display it. One column is "INSURANCE" and the values are one of five insurance categories. Print counts of how many rows are in each insurance category (hint: use `groupby()` again).

```
[4]: # YOUR CODE HERE #
admissions_df = pd.read_csv(ROOT + "/mimic_database/ADMISSIONS.csv")
display(admissions_df)

insurance_df = admissions_df.groupby(["INSURANCE"])
print(insurance_df.size())
# END CODE #
```

	ROW_ID	SUBJECT_ID	HADM_ID	ADMITTIME	DISCHTIME	\
0	21	22	165315	2196-04-09 12:26:00	2196-04-10 15:54:00	
1	22	23	152223	2153-09-03 07:15:00	2153-09-08 19:10:00	
2	23	23	124321	2157-10-18 19:34:00	2157-10-25 14:00:00	
3	24	24	161859	2139-06-06 16:14:00	2139-06-09 12:48:00	
4	25	25	129635	2160-11-02 02:06:00	2160-11-05 14:55:00	
...	...	...	...	...	...	
58971	58594	98800	191113	2131-03-30 21:13:00	2131-04-02 15:02:00	
58972	58595	98802	101071	2151-03-05 20:00:00	2151-03-06 09:10:00	
58973	58596	98805	122631	2200-09-12 07:15:00	2200-09-20 12:08:00	



58974	58597	98813	170407	2128-11-11 02:29:00	2128-12-22 13:11:00
58975	58598	98813	190264	2131-10-25 03:09:00	2131-10-26 17:44:00

	DEATHTIME	ADMISSION_TYPE	ADMISSION_LOCATION	\
0	NaN	EMERGENCY	EMERGENCY ROOM ADMIT	
1	NaN	ELECTIVE	PHYS REFERRAL/NORMAL DELI	
2	NaN	EMERGENCY	TRANSFER FROM HOSP/EXTRAM	
3	NaN	EMERGENCY	TRANSFER FROM HOSP/EXTRAM	
4	NaN	EMERGENCY	EMERGENCY ROOM ADMIT	
...	...	...	...	
58971	NaN	EMERGENCY	CLINIC REFERRAL/PREMATURE	
58972	2151-03-06 09:10:00	EMERGENCY	CLINIC REFERRAL/PREMATURE	
58973	NaN	ELECTIVE	PHYS REFERRAL/NORMAL DELI	
58974	NaN	EMERGENCY	EMERGENCY ROOM ADMIT	
58975	NaN	EMERGENCY	CLINIC REFERRAL/PREMATURE	

	DISCHARGE_LOCATION	INSURANCE	LANGUAGE	RELIGION	\
0	DISC-TRAN CANCER/CHLDRN H	Private	NaN	UNOBTAINABLE	
1	HOME HEALTH CARE	Medicare	NaN	CATHOLIC	
2	HOME HEALTH CARE	Medicare	ENGL	CATHOLIC	
3	HOME	Private	NaN	PROTESTANT QUAKER	
4	HOME	Private	NaN	UNOBTAINABLE	
...	...	...	...	...	
58971	HOME	Private	ENGL	NOT SPECIFIED	
58972	DEAD/EXPIRED	Medicare	ENGL	CATHOLIC	
58973	HOME HEALTH CARE	Private	ENGL	NOT SPECIFIED	
58974	SNF	Private	ENGL	CATHOLIC	
58975	HOME	Private	ENGL	CATHOLIC	

	MARITAL_STATUS	ETHNICITY	EDREGTIME	EDOUTTIME	\
0	MARRIED	WHITE	2196-04-09 10:06:00	2196-04-09 13:24:00	
1	MARRIED	WHITE	NaN	NaN	
2	MARRIED	WHITE	NaN	NaN	
3	SINGLE	WHITE	NaN	NaN	
4	MARRIED	WHITE	2160-11-02 01:01:00	2160-11-02 04:27:00	
...	...	...	...	...	
58971	SINGLE	WHITE	2131-03-30 19:44:00	2131-03-30 22:41:00	
58972	WIDOWED	WHITE	2151-03-05 17:23:00	2151-03-05 21:06:00	
58973	MARRIED	WHITE	NaN	NaN	
58974	MARRIED	WHITE	2128-11-10 23:48:00	2128-11-11 03:16:00	
58975	MARRIED	WHITE	2131-10-25 00:08:00	2131-10-25 04:35:00	

	DIAGNOSIS	\
0	BENZODIAZEPINE OVERDOSE	
1	CORONARY ARTERY DISEASE\CORONARY ARTERY BYPASS...	
2	BRAIN MASS	
3	INTERIOR MYOCARDIAL INFARCTION	
4	ACUTE CORONARY SYNDROME	

```

...
58971                                TRAUMA
58972                                SAH
58973                RENAL CANCER/SDA
58974                                S/P FALL
58975                INTRACRANIAL HEMORRHAGE

```

```

HOSPITAL_EXPIRE_FLAG  HAS_CHARTEVENTS_DATA
0                      0                      1
1                      0                      1
2                      0                      1
3                      0                      1
4                      0                      1
...
58971                  0                      1
58972                  1                      1
58973                  0                      1
58974                  0                      0
58975                  0                      1

```

[58976 rows x 19 columns]

```

INSURANCE
Government      1783
Medicaid        5785
Medicare        28215
Private          22582
Self Pay         611
dtype: int64

```

### Q2.3 events text database

We'll be using the raw clinician notes from `NOTEEVENTS.CSV`, also in `ROOT/mimic_database`. This is a big file, so load in just the first 10 rows, and print them. You might notice that all 'CATEGORY' columns are type 'Discharge summary'.

Then print the full text of the first row, (the 'TEXT' column). Note that this should be over 10 lines of visible text; if you don't select the entry correctly then you may see an abbreviated version.

```

[5]: # YOUR CODE HERE #
first_nrows = 10
note_df = pd.read_csv(ROOT + "/mimic_database/NOTEEVENTS.csv",
↳nrows=first_nrows)
print(note_df)
print(note_df["TEXT"].iloc[0])
# END CODE #

```

```

ROW_ID  SUBJECT_ID  HADM_ID  CHARTDATE  CHARTTIME  STORETIME  \
0      174        22532   167853   2151-08-04      NaN      NaN
1      175        13702   107527   2118-06-14      NaN      NaN

```

2	176	13702	167118	2119-05-25	NaN	NaN
3	177	13702	196489	2124-08-18	NaN	NaN
4	178	26880	135453	2162-03-25	NaN	NaN
5	179	53181	170490	2172-03-08	NaN	NaN
6	180	20646	134727	2112-12-10	NaN	NaN
7	181	42130	114236	2150-03-01	NaN	NaN
8	182	56174	163469	2118-08-12	NaN	NaN
9	183	56174	189681	2118-12-09	NaN	NaN

	CATEGORY	DESCRIPTION	CGID	ISERROR	\
0	Discharge	summary	Report	NaN	NaN
1	Discharge	summary	Report	NaN	NaN
2	Discharge	summary	Report	NaN	NaN
3	Discharge	summary	Report	NaN	NaN
4	Discharge	summary	Report	NaN	NaN
5	Discharge	summary	Report	NaN	NaN
6	Discharge	summary	Report	NaN	NaN
7	Discharge	summary	Report	NaN	NaN
8	Discharge	summary	Report	NaN	NaN
9	Discharge	summary	Report	NaN	NaN

	TEXT
0	Admission Date: [**2151-7-16**] Dischar...
1	Admission Date: [**2118-6-2**] Discharg...
2	Admission Date: [**2119-5-4**] D...
3	Admission Date: [**2124-7-21**] ...
4	Admission Date: [**2162-3-3**] D...
5	Admission Date: [**2172-3-5**] D...
6	Admission Date: [**2112-12-8**] ...
7	Admission Date: [**2150-2-25**] ...
8	Admission Date: [**2118-8-10**] ...
9	Admission Date: [**2118-12-7**] ...
Admission Date: [**2151-7-16**] Discharge Date: [**2151-8-4**]	

Service:  
ADDENDUM:

RADIOLOGIC STUDIES: Radiologic studies also included a chest CT, which confirmed cavitary lesions in the left lung apex consistent with infectious process/tuberculosis. This also moderate-sized left pleural effusion.

HEAD CT: Head CT showed no intracranial hemorrhage or mass effect, but old infarction consistent with past medical history.

ABDOMINAL CT: Abdominal CT showed lesions of

T10 and sacrum most likely secondary to osteoporosis. These can be followed by repeat imaging as an outpatient.

```
                [**First Name8 (NamePattern2) **] [**First Name4  
(NamePattern1) 1775**] [**Last Name (NamePattern1) **], M.D. [**MD Number(1)  
1776**]
```

Dictated By:**[\*\*Hospital 1807\*\*]**  
MEDQUIST36

D: **[\*\*2151-8-5\*\*]** 12:11  
T: **[\*\*2151-8-5\*\*]** 12:21  
JOB#: **[\*\*Job Number 1808\*\*]**

At the start of this assignment you ran `readmission_utils.get_notes_and_labels(ROOT)`. It did the following. - Sampled about 1000 patient admissions from `ADMISSIONS.csv` and extracted their discharge summary text from `NOTEEVENTS.csv`. - For the admission `i`: - `notes[i]` is the discharge summary text. - `labels[i]` is a 1 if that patient was readmitted after 30 days, and 0 otherwise. - The entries are randomly shuffled

Use the next code cell to compute the amount of class imbalance in the sampled dataset. You can just print the counts of labels, or show them as a histogram.

```
[6]: # YOUR CODE HERE #  
print("Label 0 count: " + str(labels_admission.count(0)))  
print("Label 1 count: " + str(labels_admission.count(1)))  
# YOUR CODE HERE #
```

Label 0 count: 502  
Label 1 count: 449

### 1.1.1 Word embeddings

#### Q2.4 tokenizers

We now want to learn a word embedding model, so that we can convert the words in `notes` to vectors that can be fed into a deep learning model.

The first step is to tokenize the notes. Execute the code in the next cell. You can read about what it's doing [here](#).

```
[7]: vocab_size = VOCAB_SIZE  
tokenizer = tf.keras.preprocessing.text.Tokenizer(  
    num_words=vocab_size,  
    oov_token="<unk>",  
    filters='!"#$%&()*+.,:;=?@[\\]^_`{|}~/\n')  
tokenizer.fit_on_texts(notes)
```

```
notes_seq = tokenizer.texts_to_sequences(notes)
```

Notice that the `tokenizer` is only going to index the 500 most common words, and set the remainder to `<unk>`. Try printing the result of `tokenizer.index_word` and `tokenizer.word_index`. (Please do not actually print these dictionaries when you submit the assignment; they print ~1000 lines of text).

Explain the content of `notes_seq`, and how it relates to `notes`.

*Written answer:* `notes_seq` transforms each note in `notes` to a sequence of integers. This list of integer sequences encodes the words in our notes, which means the words are replaced by its corresponding integer value from the `word_index` dictionary, which is the 500 most common words.

After running `tokenizer.fit_on_texts(notes)`, the `tokenizer` object stores the word counts that are in `notes`. Complete the below function to return an array of words with an array of their word counts. The arrays do not need to be sorted. Then execute the code to print the 50 most common words.

```
[8]: def get_words_and_counts(tokenizer):  
    """  
    Return an array of `words` and an array of their `counts` for the dataset_  
    ↳fitted to  
    Keras `tokenizer` object, so that words[i] appear counts[i] times. The_  
    ↳array does  
    not need to be sorted.  
  
    Parameters:  
    tokenizer (tf.keras.preprocessing.text.Tokenizer), prefitted tokenizer.  
  
    Returns:  
    vocab_words_sorted (np.array(str)) of words trained on `tokenizer`.  
    vocab_words_counts_sorted (np.array(int)) word counts so that counts[i] is_  
    ↳count of words[i]  
    """  
    # YOUR CODE HERE #  
    word_counts = tokenizer.word_counts  
    vocab_words, vocab_words_counts = list(word_counts.keys()),_  
    ↳list(word_counts.values())  
    # END CODE #  
    return vocab_words, vocab_words_counts  
  
# Provided code for printing the 50 most common words #  
n=50  
vocab_words, vocab_words_counts = get_words_and_counts(tokenizer)  
indx_sorted = np.argsort(np.array(vocab_words_counts))[:-1]  
vocab_words_sorted, vocab_words_counts_sorted = np.  
    ↳array(vocab_words)[indx_sorted], np.array(vocab_words_counts)[indx_sorted]  
print("Cnt\t Word")
```

```

for i in range(n):
    print(f"{vocab_words_counts_sorted[i]}\t{repr(vocab_words_sorted[i])}" )

```

Cnt	Word
36576	'the'
30778	'and'
26814	'to'
25842	'of'
25136	'was'
18829	'with'
17399	'a'
16993	'on'
15605	'1'
14634	'in'
13325	'for'
12865	'2'
11495	'no'
11416	'mg'
10556	'patient'
10298	'tablet'
9949	'is'
8592	'he'
8210	'blood'
8014	'po'
7915	'5'
7758	'at'
7615	'3'
7178	'name'
7031	'she'
6906	'as'
6803	'or'
6713	'discharge'
6666	'daily'
6554	'day'
6485	'4'
6361	'his'
6290	'sig'
6201	'one'
5782	'-'
5718	'history'
5438	'0'
5325	'her'
5257	'6'
5093	'left'
5081	'last'
4704	'were'
4379	's'

```
4355      'had'
4248      '7'
4247      'by'
4247      'be'
4158      '8'
4083      'admission'
4069      'right'
```

## 1.2 Word2Vec

We will now implement the Word2Vec skip-gram model. This is similar to the regular skip-gram model, but with negative sampling and subsampling (which we'll explain soon). Some background resources you may be interested in are the original paper, [Distributed Representations of Words and Phrases and their Compositionality](#), and this blog post, [Illustrated Word2Vec](#).

Here is a high level description of the Word2Vec model: - Take a 'target\_word', one one 'similar' (positive context) word, and 4 'dissimilar' (negative context) words. These words are represented as integers. - Embed each word into a vector representation (e.g. a 32-dim vector). This component is the *word embedding layer*. - Then, taking the word embeddings, predict which of the 5 context words is the 'positive context' word.

So we show the model the target word: `> [target_word]`

And 5 context words:

```
[pos_context_word, neg_context_word_1, neg_context_word_2,
neg_context_word_3, neg_context_word_4]
```

Since the positive context is at index 0, the label we train the model to predict is:

```
[1,0,0,0,0]
```

After training, we take the word embedding model and use it for other nlp tasks, like readmission prediction.

### Q2.5 poitive context words

A positive `context_word` for a `target_word` is one of the previous 2 words or the next 2 words. For example, if our sentence is:

Started on ceftriaxone and azithromycin in the ED, continued in the MICU.

And the `target_word=ceftriaxone`, then the positive context words are **started**, **on**, **and**, and **azithromycin**.

However the following are NOT positive context words because they are more than 2 words away from the target word: **ED** and **MICU**.

The idea motivating the skip-gram model is that words with similar contexts should have similar word embeddings, and we are going to enforce this when we train the Word2Vec model. Based on the examples just given of what are NOT examples of positive context words, what is one weakness of the skip-gram model for learning word embeddings?

*Written Answer:* One of the main shortcomings of the skip-gram model for learning word embeddings is that words that are more than 2 words away from the target word are not considered.

As such, some sentences that have important context words that are more than 2 words away are instead not considered and thus will affect the accuracy of our skip-gram model. In addition, skip-gram models fail to consider combined word phrases and the nuance of polysemy - for example “New York” is a single word, but the skip-gram model treats it as two separate words “New” “York”.

## Q2.6 defining skipgram contexts

We’ll build the dataset over a few functions. Note that we’re working with tokenized words from `notes_seq`, so all the data will be integers instead of strings.

In the next cell, complete the function `build_target_contexts`. You should iterate over each note, and then iterate over each word token in each note to create an array of `targets` and an array of `positive_contexts` for those targets. E.g. suppose the start of `notes_seq` is this: `> notes_seq[[1,6,3,4,7,8,6,...], [...], [...], ...]`

Then one valid data point will be `targets[i]=4` and `positive_contexts[i]=[6,3,7,8]`.

We set a 2-length context window, so any target can have between 0 and 4 positive context words (some targets will be at the start or end of the sequence and so they have fewer than 4 context words). If a target has fewer than 4 context words, then do not add it to the dataset. This is a simplification that shouldn’t affect the dataset too much since the individual notes are long.

Note also that if the word 7 appears 100 times in the text, then it will appear 100 times in `targets` as well (unless it’s omitted for having fewer than 4 context words).

The expected shape for `targets` is `(n,)`, and for `positive_contexts` is `(n,4)`. This function can run in under 20 seconds when `len(notes)<1000`.

```
[9]: def build_target_contexts(notes_seq, context_window=2):
    """
    Given a `notes_seq`, a list of lists of tokens, add each valid token to a
    numpy array `targets`, and add its positive context window to numpy array
    `positive_contexts`. The contexts are with a window `context_window`
    ↪forward
    and `context_window` back.

    All words are tokenized (represented by ints).
    E.g. for the sequence `[1,5,2,8,3,0,7]`, with context_window=2
    One returned array would be
        targets[i] = 8
        positive_contexts[i] = [5,2,3,0]

    Invalid tokens:
    In the above example, if the target is near the edges, the context vector
    ↪will be
    smaller than 4, e.g.
        targets[i] = 7
        positive_contexts[i] = [3,0]
    In this case, where len(context_window)!=2*context_window, we omit the data
    ↪point.
```



```

    Args
    notes_seq (List[List[int]]): A list of note representations, so that
    ↪notes_seq[i]
        is note i, represented by an list of token ids (which are ints).
    context_window (int): the word-distance back and forward that is still in
    ↪context.

    Returns:
    targets (np.array[int]): indices for the target words.
    positive_contexts (np.array[int,int]): array of array of context words. The
    ↪shape
        will be (n,2*context_window).
    """
    targets, positive_contexts = [], []
    for note in tqdm.tqdm(notes_seq):
        # YOUR CODE HERE #
        for i in range(len(note)):
            target = note[i]
            slc1 = slice( max(0,i-context_window), i)
            slc2 = slice( i+1, i+context_window+1)

            context = note[slc1] + note[slc2]
            if len(context) != 2*context_window:
                continue
            targets.append(target)
            positive_contexts.append(np.array(context))
        # END CODE #
    assert len(targets)==len(positive_contexts)
    return np.array(targets), np.array(positive_contexts)

# Run build_target
targets, positive_contexts = build_target_contexts(notes_seq, 2)
print(targets.shape, positive_contexts.shape, '\n')

# To verify results make sense, print the first tokens of the first note, and
↪the first set of targets and contexts #
# The targets and contexts should be the first valid ngrams of the printed note
↪#
print("Start of the dataset:")
print(notes_seq[0][:20])
print(f"\nTargets\t\tPositive contexts")
for i in range(10):
    print(f"{targets[i]}\t\t{positive_contexts[i]}")

```

100%|

| 951/951 [00:03<00:00, 244.28it/s]

```
(1542542,) (1542542, 4)
```

Start of the dataset:

```
[50, 61, 1, 29, 61, 1, 61, 5, 323, 1, 320, 358, 120, 351, 165, 1, 332, 68, 319, 266]
```

Targets	Positive contexts
1	[50 61 29 61]
29	[61 1 61 1]
61	[ 1 29 1 61]
1	[29 61 61 5]
61	[ 61 1 5 323]
5	[ 1 61 323 1]
323	[ 61 5 1 320]
1	[ 5 323 320 358]
320	[323 1 358 120]
358	[ 1 320 120 351]

## Q2.7 subsampling

In Q2.3, we saw a very high frequency of simple words like ‘the’, ‘and’, and ‘to’. One trick used in Word2Vec is ‘subsampling’; we want to sample more frequent words less often. In the below cell, we provide a function that does subsampling for you. We’ll explain how it works, and then ask a question.

The `do_subsampling` function checks the target words in the dataset, and removes words at random, but it removes frequent words with a higher probability. Here is how it works: - It create a `sampling_table` (see the keras API [see documentation](#)). It has size `VOCAB_SIZE`, so it returns a `VOCAB_SIZE`-element array containing probabilities. - The `i`th most common word should have a sampling probability of `sampling_table[i]`. For example `sampling_table[0]=0.00315` is the most common word and is sampled 0.3% of the time, while `sampling_table[-1]=0.184` is the least common word and is sampled 18% of the time. Note that these numbers depend on `sampling_factor` argument which is a chosen hyperparameter that could be tuned. - For each word, look up its sampling rate from `sampling_table`. It turns out that the Keras tokenizer indexes words in order of decreasing frequency, so the sampling rate for word `token_id` will be `sampling_table[token_id]`. - Remove words at random according to its sampling rate.

Run the code and then answer the written question.

```
[10]: ### provided code for subsampling ###
def do_subsampling(targets, positive_contexts, vocab_size=500,
    ↪sampling_factor=1e-05):
    """
    Given a list of targets and contexts output from build_target_contexts,
    ↪reduce
    the size by removing words with a probability from
    tf.keras.preprocessing.sequence.make_sampling_table.

    Args:
```

```

    targets (np.array[int]): same as output of build_target_contexts.
    positive_contexts (np.array[int,int]): same as output of
    ↪ build_target_contexts.

    Returns:
    targets_subsampled (np.array[int]): reduced version of targets after
    ↪ subsampling
    positive_contexts_subsampled (np.array[int,int]): reduced version of
    ↪ positive_contexts after subsampling
    """
    # generate sampling table
    sampling_table = tf.keras.preprocessing.sequence.
    ↪ make_sampling_table(vocab_size, sampling_factor=sampling_factor)
    # lookup sampling rates, using the fact that get sample rates
    sampling_rates = sampling_table[targets]
    # generate random numbers to compare to the sampling rates
    random_nums = np.random.sample(len(sampling_rates))
    # generate True/False for whether to keep this sample
    do_sample = random_nums < sampling_rates
    # create new array having filtered some words
    targets_subsampled = targets[do_sample]
    positive_contexts_subsampled = positive_contexts[do_sample]
    return targets_subsampled, positive_contexts_subsampled
### provided code for subsampling ###

# run subsampling
print(f"Original dataset shapes      {targets.shape}, {positive_contexts.
    ↪ shape}")
targets_subsampled, positive_contexts_subsampled = do_subsampling(targets,
    ↪ positive_contexts
    ,
    ↪ vocab_size=VOCAB_SIZE)
print(f"Dataset shapes after subsampling {targets_subsampled.shape},
    ↪ {positive_contexts_subsampled.shape}")

```

Original dataset shapes (1542542,), (1542542, 4)

Dataset shapes after subsampling (64398,), (64398, 4)

According to the [original paper](#), what are the benefits of subsampling?

*Written answer:* By subsampling frequent words, we are able to obtain significant speedup and also learn more regular word representations. As a result, we improve accuracy of representation of less frequent words and better vector representations of frequent words. As a result, we counter the imbalance between rare and infrequent words.

## Q2.8 Negative Sampling

Before Q2.4, we explained how Word2Vec works. Recall that we need to give the model a target word, a positive context word, and 4 negative context words. The model's task is to predict which

word is the positive context word.

Our next step is to generate the negative context words. Firstly we provide a function for generating negative samples. Run the next cell and look at the example usage to make sure you understand what it's doing.

```
[11]: def get_negative_samples(target, postive_context, num_ns=4, vocab_size=500):  
    """  
    Given a target word index and a list of positive context integers, randomly  
    sample new integers not in `target` or `postive_context`. Generate `num_ns`  
    ↪ samples.  
  
    Args  
    target (int): target int that should not be in `negative_context`  
    postive_context (List(int)): positive int that should not be in  
    ↪ `negative_context`  
    num_ns (int): number of negative samples to return.  
    vocab_size (int): size of vocabulary indexed by ints [0,vocab_size].  
  
    Returns:  
    negative_context (np.array[int]). Negative context tokens shape (num_ns,).  
    """  
    neg_samples_candidates = list(set(np.arange(vocab_size)) -  
    ↪ set(postive_context) - set([target]))  
    negative_context = np.random.choice(neg_samples_candidates, size=num_ns,  
    ↪ replace=False)  
    return negative_context  
  
target_test = 5  
positive_context_test = [1,2,8,9]  
vocab_size_test = 10  
num_ns_test=4  
print(f"Test generating 10 sets of negative sampling with target word  
    ↪ {target_test}, vocab_size {vocab_size_test}, positive context  
    ↪ {positive_context_test}\n")  
for i in range(10):  
    print(get_negative_samples(target_test, positive_context_test, num_ns_test,  
    ↪ vocab_size_test))
```

Test generating 10 sets of negative sampling with target word 5, vocab\_size 10,  
positive context [1, 2, 8, 9]

```
[4 3 0 6]  
[6 0 3 7]  
[0 3 4 7]  
[4 3 7 6]  
[3 4 7 0]  
[6 7 0 4]
```

```
[0 7 3 6]
[7 4 0 6]
[0 4 7 3]
[7 0 4 6]
```

We already have `targets_subsampled` and `positive_contexts_subsampled`. Let's now produce a third array `negative_contexts_subsampled` which will hold our negative samples.

We are storing each target with its entire context window: `> targets[i]=8`, and `positive_contexts[i]=[5,2,3,0]`

But in the final model we'll actually want to generate 4 training samples with this, one for each context word. So we'll get samples `[8,5]`, `[8,2]`, `[8,3]`, and `[8,0]`. And for each one of these pairs we want to generate `NUM_NS=4` negative samples. Since there are 4 training pairs in each `positive_contexts[i]`, we will need to generate `4*NUM_NS=20` negative samples for each `targets[i]`.

Implement this in the next function by making use of the function `get_negative_samples`. It should run in about 1 minute.

```
[12]: def build_target_positive_and_negative_contexts(targets_subsampled,
    ↪ positive_contexts_subsampled,
    num_ns=4, vocab_size=500):
    """
    Generate negative context words for `targets_subsampled` and
    ↪ `positive_contexts_subsampled`.
    Uses `get_negative_samples` method.

    Args:
    targets_subsampled (np.array[int]): same as output from `do_subsampling`.
    positive_contexts_subsampled (np.array([int,int])) same as output from
    ↪ `do_subsampling`.
    num_ns (int): number of negative samples per array.
    vocab_size (int): vocab size. All all_targetes[i]<vocab_size.

    Returns:
    negative_contexts_subsampled (np.array[int,int]): shape
    ↪ (n_samples,n_p*num_ns) where
        n_p is the number of context words per target,
    ↪ n_p=positive_contexts_subsampled.shape[1].
        The negative context words for the p samples.
    """
    n, n_p = positive_contexts_subsampled.shape
    negative_contexts_subsampled = np.zeros((n, n_p*num_ns))

    for i in tqdm.trange(n):
        # YOUR CODE HERE #
        t = targets_subsampled[i]
        postive_context = positive_contexts_subsampled[i]
```

```

        neg_samples = get_negative_samples(t, positive_context, n_p*num_ns, vocab_size)
        negative_contexts_subsampled[i] = neg_samples

    # END CODE #
    return negative_contexts_subsampled

negative_contexts_subsampled = \
    build_target_positive_and_negative_contexts(targets_subsampled, positive_contexts_subsampled,
                                                num_ns=NUM_NS, vocab_size=VOCAB_SIZE)
print(targets_subsampled.shape, positive_contexts_subsampled.shape, negative_contexts_subsampled.shape) # expect (n,) (n,4) (n,16)

```

```

100%|          | 64398/64398 [00:12<00:00, 5324.25it/s]
(64398,) (64398, 4) (64398, 16)

```

The previous cell explained how each row in `targets[i]` will have 4 data points: one for each positive context. In the next cell we create the final dataset with some simple reshape operations.

Look at the shape of these arrays. They have 4 times the rows as the previous cell. Each target has 1 positive context, and 4 negative contexts.

```

[13]: n, n_p = positive_contexts_subsampled.shape
dataset_targets = np.reshape(np.repeat(targets_subsampled, 4, axis=None), (4*n,1))
dataset_positive_contexts = np.reshape(positive_contexts_subsampled, (4*n,1))
dataset_negative_contexts = np.reshape(negative_contexts_subsampled, (4*n,4))
print(dataset_targets.shape, dataset_positive_contexts.shape, dataset_negative_contexts.shape)

```

```

(257592, 1) (257592, 1) (257592, 4)

```

## Q2.9 Word2Vec word embedding layer

Execute the next cell. It combines the positive and negative context arrays into one array that will be passed to Word2Vec. The model will try to predict which of the 5 samples is the positive context.

The below code also generates ground-truth labels `labels`. Since we always put the positive context word as the first element, then all labels will be `[1,0,0,0,0]`.

```

[14]: dataset_contexts = np.hstack((dataset_positive_contexts, dataset_negative_contexts))
dataset_labels = np.zeros_like(dataset_contexts)
dataset_labels[:,0] = 1

```

```

dataset = tf.data.Dataset.from_tensor_slices(((dataset_targets, ↵
↵dataset_contexts), dataset_labels))
dataset = dataset.shuffle(10000).batch(1024, drop_remainder=True)
print(dataset)

print("targets example")
print(dataset_targets[:10])
print("context example (positive context in index 0)")
print(dataset_contexts[:10])
print("labels example")
print(dataset_labels[:10])

```

```

<BatchDataset shapes: (((1024, 1), (1024, 5)), (1024, 5)), types: ((tf.int64,
tf.float64), tf.float64)>

```

targets example

```

[[323]
 [323]
 [323]
 [323]
 [332]
 [332]
 [332]
 [332]
 [135]
 [135]]

```

context example (positive context in index 0)

```

[[ 61. 204. 186. 250. 429.]
 [  5. 417. 171. 235. 175.]
 [  1. 268. 112. 189. 125.]
 [320. 479. 148.  95. 164.]
 [165.  99. 160.   2. 395.]
 [  1. 342. 337. 335. 413.]
 [ 68. 284. 462. 216. 367.]
 [319.  65. 265.  38. 411.]
 [  3. 283. 433.  72. 350.]
 [270. 473. 391. 107. 130.]]

```

labels example

```

[[1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]]

```

```

2021-10-30 02:50:46.733504: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:46.841230: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:46.842134: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:46.845823: I tensorflow/core/platform/cpu_feature_guard.cc:142]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
2021-10-30 02:50:46.846123: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:46.846896: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:46.847664: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:48.873251: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:48.874089: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:48.874900: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 02:50:48.876510: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 10819 MB memory:  -> device:
0, name: Tesla K80, pci bus id: 0000:00:04.0, compute capability: 3.7

```

We have provided most of the Word2Vec model. In the next cell you need to add the model



embedding layers (see [Keras Embedding](#) docs). We have different embedding functions. - Define `self.target_embedding` layer for the target. It expects 1-element arrays from `targets` - Define `self.context_embedding` layer for the context (positive and negative context words). It expects 5-element arrays from `targets`.

```
[15]: class Word2Vec(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, num_ns):
        super(Word2Vec, self).__init__()
        self.target_embedding = None
        self.context_embedding = None

        # YOUR CODE HERE #
        self.target_embedding = tf.keras.layers.Embedding(vocab_size,
        ↪embedding_dim, input_length=1)
        self.context_embedding = tf.keras.layers.Embedding(vocab_size,
        ↪embedding_dim, input_length=num_ns+1)
        # END CODE #

    def call(self, pair):
        target, context = pair
        target = tf.squeeze(target, axis=1)
        word_emb = self.target_embedding(target)
        # word_emb: (batch, embed)
        context_emb = self.context_embedding(context)
        # context_emb: (batch, context, embed)
        dots = tf.einsum('be,bce->bc', word_emb, context_emb)
        # dots: (batch, context)
        return dots
```

## Q2.10 training Word2Vec

Create, compile and run the model. We recommend: - 100 epochs. - 32-dim word embedding dimension. - Adam optimizer with default params. - Categorical cross entropy with the following call `tf.keras.losses.CategoricalCrossentropy(from_logits=True)`

You should get accuracy >0.9.

```
[16]: # YOUR CODE HERE #
epochs = 100
embedding_dim = 32
num_ns = 4
model_word2vec = Word2Vec(vocab_size, embedding_dim, num_ns)
model_word2vec.compile(optimizer='adam',
                        loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
                        metrics=['accuracy'])
model_word2vec.fit(dataset, epochs=epochs)
# END CODE #
```

Epoch 1/100

2021-10-30 02:50:49.678312: I  
tensorflow/compiler/mlir/mlir\_graph\_optimization\_pass.cc:185] None of the MLIR  
Optimization Passes are enabled (registered 2)

251/251 [=====] - 4s 9ms/step - loss: 1.2520 -  
accuracy: 0.6738  
Epoch 2/100  
251/251 [=====] - 2s 9ms/step - loss: 0.7752 -  
accuracy: 0.7206  
Epoch 3/100  
251/251 [=====] - 2s 9ms/step - loss: 0.7070 -  
accuracy: 0.7412  
Epoch 4/100  
251/251 [=====] - 2s 9ms/step - loss: 0.6467 -  
accuracy: 0.7657  
Epoch 5/100  
251/251 [=====] - 2s 9ms/step - loss: 0.5912 -  
accuracy: 0.7882  
Epoch 6/100  
251/251 [=====] - 2s 9ms/step - loss: 0.5462 -  
accuracy: 0.8055  
Epoch 7/100  
251/251 [=====] - 3s 10ms/step - loss: 0.5118 -  
accuracy: 0.8174  
Epoch 8/100  
251/251 [=====] - 2s 10ms/step - loss: 0.4854 -  
accuracy: 0.8266  
Epoch 9/100  
251/251 [=====] - 2s 9ms/step - loss: 0.4652 -  
accuracy: 0.8342  
Epoch 10/100  
251/251 [=====] - 2s 9ms/step - loss: 0.4488 -  
accuracy: 0.8399  
Epoch 11/100  
251/251 [=====] - 2s 9ms/step - loss: 0.4356 -  
accuracy: 0.8445  
Epoch 12/100  
251/251 [=====] - 2s 9ms/step - loss: 0.4244 -  
accuracy: 0.8481  
Epoch 13/100  
251/251 [=====] - 2s 9ms/step - loss: 0.4150 -  
accuracy: 0.8513  
Epoch 14/100  
251/251 [=====] - 2s 9ms/step - loss: 0.4069 -  
accuracy: 0.8542  
Epoch 15/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3996 -  
accuracy: 0.8569

Epoch 16/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3934 -  
accuracy: 0.8591  
Epoch 17/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3877 -  
accuracy: 0.8609  
Epoch 18/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3826 -  
accuracy: 0.8629  
Epoch 19/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3779 -  
accuracy: 0.8646  
Epoch 20/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3738 -  
accuracy: 0.8663  
Epoch 21/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3699 -  
accuracy: 0.8676  
Epoch 22/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3665 -  
accuracy: 0.8688  
Epoch 23/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3633 -  
accuracy: 0.8700  
Epoch 24/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3601 -  
accuracy: 0.8711  
Epoch 25/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3574 -  
accuracy: 0.8723  
Epoch 26/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3549 -  
accuracy: 0.8732  
Epoch 27/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3523 -  
accuracy: 0.8742  
Epoch 28/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3500 -  
accuracy: 0.8749  
Epoch 29/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3477 -  
accuracy: 0.8759  
Epoch 30/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3458 -  
accuracy: 0.8764  
Epoch 31/100  
251/251 [=====] - 3s 10ms/step - loss: 0.3438 -  
accuracy: 0.8771

Epoch 32/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3421 -  
accuracy: 0.8776  
Epoch 33/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3402 -  
accuracy: 0.8784  
Epoch 34/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3387 -  
accuracy: 0.8788  
Epoch 35/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3371 -  
accuracy: 0.8792  
Epoch 36/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3356 -  
accuracy: 0.8797  
Epoch 37/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3342 -  
accuracy: 0.8803  
Epoch 38/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3329 -  
accuracy: 0.8806  
Epoch 39/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3316 -  
accuracy: 0.8810  
Epoch 40/100  
251/251 [=====] - 2s 10ms/step - loss: 0.3304 -  
accuracy: 0.8815  
Epoch 41/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3291 -  
accuracy: 0.8818  
Epoch 42/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3280 -  
accuracy: 0.8824  
Epoch 43/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3270 -  
accuracy: 0.8826  
Epoch 44/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3259 -  
accuracy: 0.8830  
Epoch 45/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3250 -  
accuracy: 0.8833  
Epoch 46/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3240 -  
accuracy: 0.8837  
Epoch 47/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3232 -  
accuracy: 0.8839

Epoch 48/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3224 -  
accuracy: 0.8841  
Epoch 49/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3215 -  
accuracy: 0.8845  
Epoch 50/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3208 -  
accuracy: 0.8848  
Epoch 51/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3200 -  
accuracy: 0.8849  
Epoch 52/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3191 -  
accuracy: 0.8852  
Epoch 53/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3185 -  
accuracy: 0.8855  
Epoch 54/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3178 -  
accuracy: 0.8856  
Epoch 55/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3171 -  
accuracy: 0.8859  
Epoch 56/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3165 -  
accuracy: 0.8861  
Epoch 57/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3159 -  
accuracy: 0.8863  
Epoch 58/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3152 -  
accuracy: 0.8865  
Epoch 59/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3146 -  
accuracy: 0.8867  
Epoch 60/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3139 -  
accuracy: 0.8870  
Epoch 61/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3134 -  
accuracy: 0.8872  
Epoch 62/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3128 -  
accuracy: 0.8874  
Epoch 63/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3124 -  
accuracy: 0.8876

Epoch 64/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3119 -  
accuracy: 0.8877

Epoch 65/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3115 -  
accuracy: 0.8878

Epoch 66/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3109 -  
accuracy: 0.8880

Epoch 67/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3105 -  
accuracy: 0.8882

Epoch 68/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3100 -  
accuracy: 0.8883

Epoch 69/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3097 -  
accuracy: 0.8884

Epoch 70/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3091 -  
accuracy: 0.8886

Epoch 71/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3088 -  
accuracy: 0.8888

Epoch 72/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3083 -  
accuracy: 0.8887

Epoch 73/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3080 -  
accuracy: 0.8890

Epoch 74/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3076 -  
accuracy: 0.8891

Epoch 75/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3072 -  
accuracy: 0.8892

Epoch 76/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3067 -  
accuracy: 0.8894

Epoch 77/100  
251/251 [=====] - 2s 10ms/step - loss: 0.3065 -  
accuracy: 0.8894

Epoch 78/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3062 -  
accuracy: 0.8895

Epoch 79/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3058 -  
accuracy: 0.8897

Epoch 80/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3055 -  
accuracy: 0.8897  
Epoch 81/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3052 -  
accuracy: 0.8899  
Epoch 82/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3048 -  
accuracy: 0.8902  
Epoch 83/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3044 -  
accuracy: 0.8902  
Epoch 84/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3041 -  
accuracy: 0.8904  
Epoch 85/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3040 -  
accuracy: 0.8905  
Epoch 86/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3036 -  
accuracy: 0.8906  
Epoch 87/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3033 -  
accuracy: 0.8907  
Epoch 88/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3030 -  
accuracy: 0.8909  
Epoch 89/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3027 -  
accuracy: 0.8909  
Epoch 90/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3023 -  
accuracy: 0.8910  
Epoch 91/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3022 -  
accuracy: 0.8910  
Epoch 92/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3019 -  
accuracy: 0.8911  
Epoch 93/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3016 -  
accuracy: 0.8912  
Epoch 94/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3014 -  
accuracy: 0.8913  
Epoch 95/100  
251/251 [=====] - 2s 9ms/step - loss: 0.3011 -  
accuracy: 0.8915

```

Epoch 96/100
251/251 [=====] - 2s 9ms/step - loss: 0.3009 -
accuracy: 0.8915
Epoch 97/100
251/251 [=====] - 2s 9ms/step - loss: 0.3006 -
accuracy: 0.8917
Epoch 98/100
251/251 [=====] - 2s 9ms/step - loss: 0.3005 -
accuracy: 0.8917
Epoch 99/100
251/251 [=====] - 2s 9ms/step - loss: 0.3002 -
accuracy: 0.8919
Epoch 100/100
251/251 [=====] - 2s 9ms/step - loss: 0.2999 -
accuracy: 0.8920

```

[16]: <keras.callbacks.History at 0x7f6606b014d0>

What is the baseline accuracy for this prediction task? In other words, what ‘accuracy’ would you expect if we were randomly guessing predictions.

*Written answer:* I would expect an accuracy of 20% because there is a 1/5 chance that you randomly guess the positive label correctly

## Q2.11 word embeddings

Word2Vec learns to predict positive-context words from a list of positive- and negative-context words. In order to do this, Word2Vec must embed integer tokens into fixed-length vectors called ‘word embeddings’. The point of doing Word2Vec is to get these embeddings, and use them for downstream prediction tasks.

Complete the following function to get a matrix that will store the word embeddings for our vocabulary. You should use the `target_embedding` layer.

```

[17]: def get_word_embeddings(model_word2vec, vocab_size):
    """
    Take the target word embedding layer from `model_word2vec`. Produce an
    →embedding
    vector for a vocabulary with size vocab_size, so that `embeddings[i]`
    →returns the
    word embedding vector for the ith word.

    Args:
    model_word2vec (class Word2Vec): a trained Word2Vec model.
    vocab_size (int): vocab size; the model must have been trained on this size.

    Returns:
    embedding (np.array[float,float]): with shape (vocab_size, embedding_dim)
    """
    embeddings = None

```



```

# YOUR CODE HERE #
embeddings = np.asarray(model_word2vec.target_embedding.get_weights())
embeddings = np.squeeze(embeddings, axis=0)
# END CODE #
return embeddings

embeddings = get_word_embeddings(model_word2vec, VOCAB_SIZE)
print(embeddings.shape)    # expect (VOCAB_SIZE, embedding_dim)

```

(500, 32)

## Q2.12 nearest words

Word2Vec is trained so that words with similar contexts have similar word embeddings (as measured by cosine similarity).

We provide the function `find_nearest_words` below. Given a target word, it returns a string of the nearest words in the embedding space.

All you have to do for this question is add some words to the `chosen_words`, e.g. ['bleeding', 'pain', 'and'], and then execute the code in the cell.

```

[18]: from sklearn.metrics.pairwise import cosine_similarity
dists = cosine_similarity(embeddings, embeddings)

def find_nearest_words(word, embeddings, tokenizer):
    """
    Given
    Args:
    word (str): target word.
    embeddings (np.array[float,float]): same as output to get_word_embeddings.
    tokenizer: (tf.keras.preprocessing.text.Tokenizer) with vocabulary
    →corresponding
        to `embeddings` st tokenizer.word_index[word]=i is the ith column of
    →`embeddings`.
    """
    dists = cosine_similarity(embeddings, embeddings)
    idx = tokenizer.word_index.get(word,501)
    if idx>=VOCAB_SIZE:
        return 'ERROR: NOT IN VOCAB'
    nearest = np.argsort(dists[idx])[:-1]
    nearest_words = ''
    for j in range(1,15):
        nearest_words += tokenizer.index_word[nearest[j]] + ', '
    return nearest_words

chosen_words = None
# YOUR CODE HERE #

```

```

chosen_words = ['bleeding', 'pain', 'and', 'off', 'breath', 'air',
↳ 'rhythm', 'started', 'medicine', 'infection']
# END CODE #

for word in chosen_words:
    nearest_words = find_nearest_words(word, embeddings, tokenizer)
    print(f"TARGET: {word}\nNEAREST: {nearest_words}")
    print("\n")

```

TARGET: bleeding  
 NEAREST: bleed, infection, gi, been, any, significant, report, fevers, some, prior, signs, not, while, symptoms,

TARGET: pain  
 NEAREST: fevers, fever, nausea, breath, vomiting, cough, abdominal, shortness, prn, constipation, chest, back, or, lower,

TARGET: and  
 NEAREST: noted, tube, her, however, continued, pleural, so, chest, <unk>, was, upper, fluid, low, which,

TARGET: off  
 NEAREST: then, discontinued, so, coumadin, which, morning, every, hours, this, placement, days, placed, q6h, patient,

TARGET: breath  
 NEAREST: pain, shortness, chest, some, without, nausea, fever, upper, lungs, cough, any, signs, bleeding, due,

TARGET: air  
 NEAREST: wall, sounds, micu, oxygen, ra, room, ct, fluid, size, 90, rhythm, in, 100, the,

TARGET: rhythm  
 NEAREST: sinus, oxygen, sounds, rate, alert, ekg, size, with, improved, decreased, echo, in, good, regular,

TARGET: started  
 NEAREST: when, discontinued, given, initially, placed, recommended, received, treated, continued, intubated, held, vancomycin, so, followed,

TARGET: medicine

NEAREST: service, sex, m, therapy, inr, after, pt, f, general, need, surgery, hematocrit, needed, facility,

TARGET: infection

NEAREST: bleeding, all, care, been, acute, time, urine, disease, bleed, level, symptoms, no, distress, antibiotics,

### 1.2.1 Prediction with word embeddings

#### Q2.13 data representation for notes

Now that we have word embeddings for our notes, lets make predictions. We will provide data-generating code, and you will define the model.

The original dataset is two equally-sized lists, so that `notes[i]` is the discharge summary of visit `i`, and `labels_admission` is a 1 if there was a readmission within 30 days, and 0 otherwise. The next cell creates the input to a Keras sequence model (`batch_size`, `n_tokens`, `embedding_dim`).

Each sequence of words can only be `n_tokens` long. Here we choose `n_tokens=512`. But the notes can be any number of tokens long. There are many strategies for choosing which word tokens to include in the note representation. We will just take the first 512 tokens from each note.

```
[19]: def convert_notes_seq_to_embeddings(notes_seq, embeddings):
    notes_word_embeddings = []
    for i, note_seq in enumerate(notes_seq):
        note_word_embeddings = tf.gather(embeddings, indices=note_seq)
        notes_word_embeddings.append(np.array(note_word_embeddings))
    return notes_word_embeddings

notes_word_embeddings = convert_notes_seq_to_embeddings(notes_seq, embeddings)

notes_first_512_words = np.zeros((len(notes), 512, embeddings.shape[-1]))
for i in range(len(notes)):
    bs = min(512, len(notes_word_embeddings[i]))
    notes_first_512_words[i, :bs] = notes_word_embeddings[i][:bs]

print(notes_first_512_words.shape) # expect (len(notes), 512, embedding_dim).
```

(951, 512, 32)

We chose to just use the first 512 word embeddings from each note as its representation. Suggest two other strategies we could have used,

*Written answer:* 1. We can use the last 512 word embeddings from each note as its representation, as this might be where the clinician includes the most relevant information of the user's condition

and has the highest information value in the note. 2. We can sample the 512 most frequent words as that is most likely the most important words and has highest information value in the note. If some words occur many times in comparison to others, it might be a good indicator that those words are particularly informative and important to the patient's health.

Then execute the next cell which create train/val/test splits

```
[20]: from sklearn.model_selection import train_test_split
DATA = notes_first_512_words
LABELS = labels_admission
X_train, X_test, y_train, y_test = train_test_split(notes_first_512_words, np.
    ↳ array(labels_admission), test_size=0.2, random_state=1)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.
    ↳ 25, random_state=1)
print("Train ", X_train.shape, y_train.shape)
print("Val   ", X_val.shape, y_val.shape)
print("Test  ", X_test.shape, y_test.shape)
```

```
Train (570, 512, 32) (570,)
Val   (190, 512, 32) (190,)
Test  (191, 512, 32) (191,)
```

## Q2.14 training

Create a prediction model including: - 1 masking layer. - 1 LSTM layer that returns a single vector (instead of a sequence of vectors). - 1 dropout layer. - 1 dense layer whose output is a prediction.

```
[23]: num_timesteps, num_features = X_train.shape[-2:]
num_lstm_units=32
# YOUR CODE HERE #
model_lstm = tf.keras.Sequential(
    layers=[
        tf.keras.layers.Masking(mask_value=0.,
                                input_shape=(num_timesteps, num_features)),
        tf.keras.layers.LSTM(num_lstm_units),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(1, activation='sigmoid')
    ]
)
# END CODE #
```

In part 1, the LSTM layer returned a value for every element in the sequence. In this problem the LSTM layer returns only the last element. Explain why this task is different.

*Written answer:* The LSTM layer returns only the last element in this problem, because it is a “many-to-one” problem. We are interested in predicting a single vector rather than a sequence of vectors as we want to predict which is the positive label. As a result, in contrast to the previous question, we only return the last element.

Compile the model with Adam, a suitable loss function, and the ‘accuracy’ metric. Train the model

for 15 epoch.

Note that this is a very difficult model to train with the dataset that we have. Your results should show decreasing loss on the train set, but you may not see any improvement in the validation loss.

```
[24]: # YOUR CODE HERE #
model_lstm.compile(loss=tf.keras.losses.BinaryCrossentropy(), optimizer=tf.
    ↳keras.optimizers.Adam(learning_rate=0.0001), metrics=['accuracy'])
epochs = 15
hist = model_lstm.fit(x=X_train, y=y_train, epochs=epochs,
    ↳validation_data=(X_val, y_val))
# END CODE #
```

Epoch 1/15

18/18 [=====] - 6s 156ms/step - loss: 0.7174 -  
accuracy: 0.5053 - val\_loss: 0.7143 - val\_accuracy: 0.4684

Epoch 2/15

18/18 [=====] - 2s 89ms/step - loss: 0.7126 - accuracy:  
0.5053 - val\_loss: 0.7065 - val\_accuracy: 0.4789

Epoch 3/15

18/18 [=====] - 2s 87ms/step - loss: 0.7149 - accuracy:  
0.4544 - val\_loss: 0.7032 - val\_accuracy: 0.4737

Epoch 4/15

18/18 [=====] - 2s 87ms/step - loss: 0.7037 - accuracy:  
0.5193 - val\_loss: 0.7005 - val\_accuracy: 0.4947

Epoch 5/15

18/18 [=====] - 2s 87ms/step - loss: 0.7082 - accuracy:  
0.4825 - val\_loss: 0.6983 - val\_accuracy: 0.5053

Epoch 6/15

18/18 [=====] - 2s 88ms/step - loss: 0.6913 - accuracy:  
0.5404 - val\_loss: 0.6972 - val\_accuracy: 0.5053

Epoch 7/15

18/18 [=====] - 2s 87ms/step - loss: 0.7112 - accuracy:  
0.5070 - val\_loss: 0.6969 - val\_accuracy: 0.5105

Epoch 8/15

18/18 [=====] - 2s 87ms/step - loss: 0.7018 - accuracy:  
0.5228 - val\_loss: 0.6962 - val\_accuracy: 0.5000

Epoch 9/15

18/18 [=====] - 2s 88ms/step - loss: 0.7006 - accuracy:  
0.5158 - val\_loss: 0.6959 - val\_accuracy: 0.5263

Epoch 10/15

18/18 [=====] - 2s 88ms/step - loss: 0.6969 - accuracy:  
0.5281 - val\_loss: 0.6953 - val\_accuracy: 0.5211

Epoch 11/15

18/18 [=====] - 2s 87ms/step - loss: 0.7024 - accuracy:  
0.5000 - val\_loss: 0.6954 - val\_accuracy: 0.5263

Epoch 12/15

18/18 [=====] - 2s 88ms/step - loss: 0.7055 - accuracy:

```
0.4930 - val_loss: 0.6953 - val_accuracy: 0.5263
Epoch 13/15
18/18 [=====] - 2s 87ms/step - loss: 0.7098 - accuracy:
0.4789 - val_loss: 0.6951 - val_accuracy: 0.5263
Epoch 14/15
18/18 [=====] - 2s 88ms/step - loss: 0.6899 - accuracy:
0.5368 - val_loss: 0.6951 - val_accuracy: 0.5211
Epoch 15/15
18/18 [=====] - 2s 88ms/step - loss: 0.6946 - accuracy:
0.5158 - val_loss: 0.6948 - val_accuracy: 0.5263
```

### **Q2.13**

Suggest 2 reasons why the validation results were poor when we trained this model on this dataset.

*Written answer:* 1. The validation set may have words we have not yet learned embeddings for. As a result, the model cannot accurately predict on this new dataset, since we don't have the learned values to predict on - this may result in poor predictive efficiency on the validation test set. 2. Inter-clinician variability may also be cause for the poor validation result. As we know, doctors differ greatly in how they write and append their patient notes, which may have accounted for the poor predictive quality on the validation test set.

### **1.3**

# A2\_part3\_clinical\_BERT\_embeddings\_and\_readmission\_prediction

October 30, 2021

## 1 Assignment 2 - part 3 - BERT Embeddings For Prediction

In this part of the assignment, we will attempt the same prediction task as part 2, but with two differences.

**Different subsequencing strategy** Models for sequence data need fixed sequence lengths. In part 2 we just used just the first ~500 words of each note. In part 3 we will break each note into 500-word chunks and train the model to classify each chunk separately. Then we will combine the chunked predictions into one prediction for the whole note. This is sometimes referred to as a ‘sliding window’ or ‘binning’. ([Here](#) is a discussion of strategies for long-text modeling with BERT.)

**Different embedding strategy** We need to convert sequences of word tokens to a vector representation that we can then use in a prediction model. In part 2 we converted each of the first 500 words into 500 Word2Vec embedding vectors, and then passed that sequence of 500 vectors to an LSTM prediction model. In part 3 we will instead convert each note sequence to a single vector. This vector is something we can get from [BERT](#), a popular transformer model. Specifically we will be using a BERT model trained on biomedical and clinical data, similar to the [ClinicalBert paper](#).

In the next cell replace ROOT with your path.

```
[1]: import readmission_utils
import tensorflow as tf
import pandas as pd
import random
import pickle
import numpy as np
import matplotlib.pyplot as plt
import bert_utils

ROOT = "/home/marchuo/assign2/" # Put your root path here"
tf.keras.backend.set_floatx('float32')
```

### 1.1 Preprocessing text data and visualization

Execute the code in the next cell, which will take about 60mins the first time you run it. It will save its results to a file in ROOT/saved\_data/texts\_to\_labels\_5000.pkl. This is the same code as part 2, except now we have 5000 notes instead of 1000.

If the file already exists then calling the function will just load the results. We also break the notes and labels into train/val/test sets.

```
[2]: notes, labels = readmission_utils.get_notes_and_labels(ROOT, 5000)
```

Found file /home/marchuo/assign2/saved\_data/texts\_to\_labels\_5000.pkl, loading

Run the following code which loads a pretrained Bert model from the [HuggingFace transformers library](#). This library provides a standard interface for tokenizers and transformers in Tensorflow and PyTorch. HuggingFace also provides a platform for researchers to share pretrained models. For example we are using [this BERT model + tokenizer](#) that has been trained on a dataset of biomedical texts.

This code should take less than 1 minute to run.

```
[3]: from transformers import AutoTokenizer, TFAutoModel
import readmission_utils
hf_model = "cambridgeltl/SapBERT-from-PubMedBERT-fulltext"
if 'tokenizer' not in locals().keys(): tokenizer = AutoTokenizer.
    ↳from_pretrained(hf_model)
if 'bert_auto_model' not in locals().keys(): bert_model = TFAutoModel.
    ↳from_pretrained(hf_model)
tokenizer = AutoTokenizer.from_pretrained(hf_model)
bert_model = TFAutoModel.from_pretrained(hf_model)
```

```
2021-10-30 03:36:39.101694: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.110661: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.111448: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.113495: I tensorflow/core/platform/cpu_feature_guard.cc:142]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
2021-10-30 03:36:39.113818: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.114531: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
```



```

read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.115303: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.508666: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.509580: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.510356: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2021-10-30 03:36:39.511049: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 10819 MB memory: -> device:
0, name: Tesla K80, pci bus id: 0000:00:04.0, compute capability: 3.7
All model checkpoint layers were used when initializing TFBertModel.

```

All the layers of TFBertModel were initialized from the model checkpoint at cambridgelt1/SapBERT-from-PubMedBERT-fulltext.  
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.  
All model checkpoint layers were used when initializing TFBertModel.

All the layers of TFBertModel were initialized from the model checkpoint at cambridgelt1/SapBERT-from-PubMedBERT-fulltext.  
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.

Now run the following data preparation code. We'll explain what it does later. It will take about ~20mins the first time it's run and it will save data to {ROOT}/saved\_data/bert\_datasets.pkl. For later runs, it will just load this file.

```
[4]: data = bert_utils.prepare_bert_datasets(ROOT, notes, labels, bert_model, ↵
    ↪tokenizer)
```

File /home/marchuo/assign2//saved\_data/bert\_datasets.pkl exists.  
Loading it

### Q3.1 BERT architecture

LSTMs model sequence dependencies using recurrence; they are recurrent neural networks or RNNs. In RNNs we pass elements of a sequence through the model one a time (sequentially). Each pass

through the RNN updates an internal state vector. Future passes through the RNN are a function of the state vector. This is how RNNs can model dependencies between elements in a sequence.

On the other hand, Bert has a transformer architecture. Transformers are state-of-the-art in most standard tasks in language modelling. Instead of processing sequence data one-at-a-time, transformers process entire sequences at once. But they still model dependencies between sequence elements. Briefly describe the mechanism that transformers use to model sequence dependencies. (You can refer to the lecture slides, or the major transformers paper, [Attention is all you need](#)).

*Written Answer:* Transformers follow the overall architecture of neural sequence transduction models, where they have an encoder-decoder structure. The encoder maps an input sequence of symbol representations to a sequence of continuous representations. Then, the decoder takes this sequence of continuous representations and generates an output sequence of symbols one at a time. The encoder-decoder structure is auto-regressive, consuming previously generated symbols. The transformer uses fully stacked self-attention and point wise layers.

The encoder has 6 layers with 2 sublayers between each - one is a multi-head self-attention mechanism and the other is a fully connected position-wise feed-forward network. The decoder has 6 layers with 1 sublayer in addition to the two aforementioned sublayers (total 3), which performs multi-head attention. The self-attention layer is masked by offsetting by 1 position and preventing positions from attending to subsequent positions, so that “the predictions for position i can depend only on the known outputs at positions less than i”.

Multi-head attention allows the model to jointly attend to information from different representation subspaces at different positions. Thus, every position in the decoder is able to attend over all positions in the input sequence. Similarly, each position in the encoder can attend to all positions in the previous layer of the encoder. Finally, every position in the decoder is able to attend over all positions in the decoder sequences up until that point.

### Q3.2 BERT pretraining

Briefly describe the 2 pretraining tasks discussed in the introduction to [the BERT paper](#).

*Written Answer:* 1) Masked LM. The researchers masked some percentage of the input tokens and predicted those masked tokens. The “final hidden vectors corresponding to the mask tokens are fed into an output softmax over the vocabulary”. In this particular paper, they masked 15% of the WordPiece tokens in each sequence at random, and they predicted the missing words. This allows them to receive a bidirectional pre-trained model. They also account for mismatch between pre-training and finetuning by varying the replacements of the tokens.

- 2) Next Sentence Prediction (NSP). The researchers pre-trained for a binarized next sentence prediction task. When choosing sentences A and B for each pre-training example, 50% of the time B is the sentence that follows A or 50% of the time a random sentence is the sentence that follows A. This is especially beneficial for Question Answering and Natural Language Inference.

### Q3.3 datasets for BERT pretraining

What is the benefit of using a BERT model that has been pretrained on biomedical text compared with, for example, a BERT model trained on Wikipedia?

*Written Answer:* By using a BERT model that has been pretrained on biomedical text, the model has been trained on word embeddings that are specifically referenced on clinician notes and is

incredibly relevant to our prediction task, which is biomedical in nature. On the other hand, if we were to use a BERT model train on Wikipedia, there would be a myriad of words that are never referenced and seen in the notes corpus, and the model would also be much less likely to have biomedical terms referenced in our clinician notes. There would be an abundance of unused words if trained on Wikipedia

### Q3.4 data chunking strategy

Let's look at some of the data we created earlier when we ran `bert_utils.prepare_bert_datasets`. First we did a train/val/test split for the notes and labels. All these variables have the suffix, `_FULL`, indicating that this is the full note, before chunking.

```
[5]: [
    train_notes_FULL, train_labels_FULL,
    val_notes_FULL, val_labels_FULL,
    test_notes_FULL, test_labels_FULL
] = data['FULL']

all_note_lengths = [len(train_notes_FULL), len(test_notes_FULL),
    ↪len(test_labels_FULL)]
print(f"train/val/test lengths {all_note_lengths} \n")
print(f"Which sum to {sum(all_note_lengths)}")
print(f"Original notes len {len(notes)}")
```

```
train/val/test lengths [2853, 951, 951]
```

```
Which sum to 4755
```

```
Original notes len 4755
```

Now we do chunking for the train, val and test sets separately (this is what we described in the “Different subsequencing strategy” section at the start of the assignment).

Take the train set for example. We break the notes into ~500 word chunks, `train_notes_CHUNKS`. We copy the labels into `train_labels_CHUNKS`. Finally, `train_idxes_CHUNKS` tells you which FULL note this CHUNK is from. Suppose `train_idxes_CHUNKS[30]=6`; this means `train_notes_CHUNKS[30]` is a subsequence of the note `train_notes_FULL[6]`.

Here is an example: - If `train_notes_FULL[0]` is about 1200 words long, then we create 3 note-chunks that will be in `train_notes_CHUNKS[0:3]` - If the label is `train_notes_FULL[0]=1` then we copy that label for each note-chunk, so `train_labels_CHUNKS[0:3]=1`. - Since these chunks are all subsequences of `train_notes_FULL[0]`, we set `train_idxes_CHUNKS[0:3]=0`.

We print the labels and idxes for the first 25 entries. You should verify that the results match your understanding of this dataset.

```
[6]: [
    train_notes_CHUNKS, train_labels_CHUNKS, train_idxes_CHUNKS,
    val_notes_CHUNKS, val_labels_CHUNKS, val_idxes_CHUNKS,
    test_notes_CHUNKS, test_labels_CHUNKS, test_idxes_CHUNKS,
] = data['CHUNKS_DATA']
```

```
print("First 20 chunks:")
print(f"Labels      : {train_labels_CHUNKS[:25]}")
print(f"Indexes.    : {train_idxs_CHUNKS[:25]}")
```

First 20 chunks:

```
Labels      : [1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0,
0, 0, 0]
Indexes.    : [0 0 0 0 0 0 1 1 1 2 2 2 2 2 2 2 2 2 3 3 3 4 4 4 4]
```

Briefly discuss the pros and cons of this data chunking strategy compared to the truncation strategy used in part 2.

*Written answer:* One of the strengths of this data chunking strategy rather than truncation is that we have more context for the note because you get bert embeddings for the whole note rather than just target words as compared to part 2. As a result, we can get a more holistic overview of the note rather than just the truncated words. Some of the cons could be that this process is more computationally expensive and time-consuming to accomplish in comparison to the truncated strategy. In addition, we'll have more imbalance between very frequent and common words in comparison to rare words, which may affect model accuracy.

### Q3.5 BERT embeddings

Finally we took these chunked notes and put them into BERT pooled embeddings.

```
[7]: [
    train_bert_pool_embeddings_CHUNKS,
    val_bert_pool_embeddings_CHUNKS,
    test_bert_pool_embeddings_CHUNKS,
] = data['CHUNKS_EMEDDINGS']
print(f"Length of train_notes_CHUNKS      {len(train_notes_CHUNKS)}")
print(f"Shape of train_bert_pool_embeddings_CHUNKS_
→ {train_bert_pool_embeddings_CHUNKS.shape}")
```

```
Length of train_notes_CHUNKS      12253
Shape of train_bert_pool_embeddings_CHUNKS (12253, 768)
```

## 2 For our prediction task:

- The x-data is `train_bert_pool_embeddings_CHUNKS`.
- They y-labels are `train_labels_CHUNKS`.

Look at the shape of `train_bert_pool_embeddings_CHUNKS` printed in the above cell. There is one single BERT embedding vector for each note chunk. This is called the “pooled BERT embedding”, and is also the  $h_{CLS}$  token output discussed in lecture. How is this embedding different to the embeddings used in part 2? Specifically talk about the shape of the data that we will pass into a prediction model.

*Written answer:* In part 2, we had a 32 dimension embedding, while we now have embeddings of size 768. There is one single BERT embedding vector for each note chunk, while the previous part

embeds each word into a vector representation of size 32.

### 2.0.1 Prediction model

**Q3.6 build and run prediction model** The inputs to our model are single-vector BERT embeddings. These embeddings should do a very good job of summarising the text, such that our prediction model can be extremely simple: - The input is the BERT embedding vector. - We have one Dense layer with 1 node output and sigmoid activation (no hidden layers).

Compile this model with. - Adam. - Binary cross entropy loss. - Metrics for accuracy and [AUC](#).

Train it for 100 epoch with batch size 128, and pass in the validation dataset.

(Optional: you can experiment with adding extra dense layers and dropout. See if you can avoid overfitting.)

```
[19]: train_x = train_bert_pool_embeddings_CHUNKS
train_y = np.array(train_labels_CHUNKS)
val_x = val_bert_pool_embeddings_CHUNKS
val_y = np.array(val_labels_CHUNKS)
test_x = test_bert_pool_embeddings_CHUNKS
test_y = np.array(test_labels_CHUNKS)

# YOUR CODE HERE #
model = tf.keras.Sequential(
    layers=[
        tf.keras.layers.Dense(1, activation='sigmoid')
    ]
)
model.compile(loss=tf.keras.losses.BinaryCrossentropy(), optimizer=tf.keras.
    ↳optimizers.Adam(learning_rate=0.0001), metrics=['accuracy', tf.keras.metrics.
    ↳AUC()])
epochs = 100
batch_size = 128
hist = model.fit(x=train_x, y=train_y, epochs=epochs, validation_data=(val_x,
    ↳val_y), batch_size=batch_size)
# END CODE #
```

Epoch 1/100

96/96 [=====] - 1s 7ms/step - loss: 0.7046 - accuracy:  
0.5085 - auc\_2: 0.5062 - val\_loss: 0.7047 - val\_accuracy: 0.5038 - val\_auc\_2:  
0.5024

Epoch 2/100

96/96 [=====] - 0s 5ms/step - loss: 0.6989 - accuracy:  
0.5139 - auc\_2: 0.5154 - val\_loss: 0.6983 - val\_accuracy: 0.5128 - val\_auc\_2:  
0.5161

Epoch 3/100

96/96 [=====] - 0s 5ms/step - loss: 0.6949 - accuracy:  
0.5247 - auc\_2: 0.5267 - val\_loss: 0.6947 - val\_accuracy: 0.5227 - val\_auc\_2:  
0.5311

Epoch 4/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6918 - accuracy: 0.5360 - auc\_2: 0.5400 - val\_loss: 0.6914 - val\_accuracy: 0.5315 - val\_auc\_2: 0.5444

Epoch 5/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6893 - accuracy: 0.5399 - auc\_2: 0.5498 - val\_loss: 0.6890 - val\_accuracy: 0.5405 - val\_auc\_2: 0.5578

Epoch 6/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6873 - accuracy: 0.5417 - auc\_2: 0.5594 - val\_loss: 0.6864 - val\_accuracy: 0.5521 - val\_auc\_2: 0.5690

Epoch 7/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6855 - accuracy: 0.5470 - auc\_2: 0.5675 - val\_loss: 0.6846 - val\_accuracy: 0.5582 - val\_auc\_2: 0.5779

Epoch 8/100  
96/96 [=====] - 1s 6ms/step - loss: 0.6836 - accuracy: 0.5487 - auc\_2: 0.5745 - val\_loss: 0.6826 - val\_accuracy: 0.5631 - val\_auc\_2: 0.5853

Epoch 9/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6823 - accuracy: 0.5550 - auc\_2: 0.5800 - val\_loss: 0.6818 - val\_accuracy: 0.5638 - val\_auc\_2: 0.5917

Epoch 10/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6810 - accuracy: 0.5566 - auc\_2: 0.5850 - val\_loss: 0.6805 - val\_accuracy: 0.5674 - val\_auc\_2: 0.5975

Epoch 11/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6798 - accuracy: 0.5606 - auc\_2: 0.5897 - val\_loss: 0.6800 - val\_accuracy: 0.5648 - val\_auc\_2: 0.6019

Epoch 12/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6792 - accuracy: 0.5634 - auc\_2: 0.5916 - val\_loss: 0.6781 - val\_accuracy: 0.5721 - val\_auc\_2: 0.6054

Epoch 13/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6781 - accuracy: 0.5652 - auc\_2: 0.5952 - val\_loss: 0.6764 - val\_accuracy: 0.5806 - val\_auc\_2: 0.6090

Epoch 14/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6777 - accuracy: 0.5648 - auc\_2: 0.5955 - val\_loss: 0.6767 - val\_accuracy: 0.5762 - val\_auc\_2: 0.6115

Epoch 15/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6767 - accuracy: 0.5683 - auc\_2: 0.5998 - val\_loss: 0.6747 - val\_accuracy: 0.5818 - val\_auc\_2: 0.6141

Epoch 16/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6764 - accuracy: 0.5715 - auc\_2: 0.6004 - val\_loss: 0.6752 - val\_accuracy: 0.5823 - val\_auc\_2: 0.6160

Epoch 17/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6756 - accuracy: 0.5710 - auc\_2: 0.6021 - val\_loss: 0.6744 - val\_accuracy: 0.5830 - val\_auc\_2: 0.6177

Epoch 18/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6746 - accuracy: 0.5737 - auc\_2: 0.6062 - val\_loss: 0.6729 - val\_accuracy: 0.5917 - val\_auc\_2: 0.6196

Epoch 19/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6745 - accuracy: 0.5745 - auc\_2: 0.6056 - val\_loss: 0.6733 - val\_accuracy: 0.5852 - val\_auc\_2: 0.6207

Epoch 20/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6740 - accuracy: 0.5773 - auc\_2: 0.6076 - val\_loss: 0.6725 - val\_accuracy: 0.5876 - val\_auc\_2: 0.6220

Epoch 21/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6736 - accuracy: 0.5755 - auc\_2: 0.6086 - val\_loss: 0.6717 - val\_accuracy: 0.5920 - val\_auc\_2: 0.6236

Epoch 22/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6731 - accuracy: 0.5765 - auc\_2: 0.6100 - val\_loss: 0.6721 - val\_accuracy: 0.5883 - val\_auc\_2: 0.6243

Epoch 23/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6732 - accuracy: 0.5766 - auc\_2: 0.6092 - val\_loss: 0.6708 - val\_accuracy: 0.5951 - val\_auc\_2: 0.6254

Epoch 24/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6724 - accuracy: 0.5772 - auc\_2: 0.6114 - val\_loss: 0.6714 - val\_accuracy: 0.5888 - val\_auc\_2: 0.6259

Epoch 25/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6720 - accuracy: 0.5783 - auc\_2: 0.6127 - val\_loss: 0.6719 - val\_accuracy: 0.5866 - val\_auc\_2: 0.6272

Epoch 26/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6718 - accuracy: 0.5817 - auc\_2: 0.6134 - val\_loss: 0.6711 - val\_accuracy: 0.5903 - val\_auc\_2: 0.6280

Epoch 27/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6714 - accuracy: 0.5794 - auc\_2: 0.6144 - val\_loss: 0.6713 - val\_accuracy: 0.5900 - val\_auc\_2: 0.6283

Epoch 28/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6712 - accuracy: 0.5791 - auc\_2: 0.6151 - val\_loss: 0.6695 - val\_accuracy: 0.5966 - val\_auc\_2: 0.6290

Epoch 29/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6706 - accuracy: 0.5810 - auc\_2: 0.6166 - val\_loss: 0.6694 - val\_accuracy: 0.5949 - val\_auc\_2: 0.6301

Epoch 30/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6705 - accuracy: 0.5784 - auc\_2: 0.6170 - val\_loss: 0.6688 - val\_accuracy: 0.5988 - val\_auc\_2: 0.6304

Epoch 31/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6702 - accuracy: 0.5820 - auc\_2: 0.6178 - val\_loss: 0.6686 - val\_accuracy: 0.5966 - val\_auc\_2: 0.6309

Epoch 32/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6698 - accuracy: 0.5837 - auc\_2: 0.6187 - val\_loss: 0.6686 - val\_accuracy: 0.5990 - val\_auc\_2: 0.6313

Epoch 33/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6696 - accuracy: 0.5821 - auc\_2: 0.6193 - val\_loss: 0.6690 - val\_accuracy: 0.5932 - val\_auc\_2: 0.6316

Epoch 34/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6694 - accuracy: 0.5806 - auc\_2: 0.6195 - val\_loss: 0.6695 - val\_accuracy: 0.5988 - val\_auc\_2: 0.6319

Epoch 35/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6692 - accuracy: 0.5829 - auc\_2: 0.6199 - val\_loss: 0.6682 - val\_accuracy: 0.5966 - val\_auc\_2: 0.6325

Epoch 36/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6688 - accuracy: 0.5825 - auc\_2: 0.6215 - val\_loss: 0.6684 - val\_accuracy: 0.5983 - val\_auc\_2: 0.6328

Epoch 37/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6685 - accuracy: 0.5852 - auc\_2: 0.6224 - val\_loss: 0.6676 - val\_accuracy: 0.5985 - val\_auc\_2: 0.6330

Epoch 38/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6682 - accuracy: 0.5859 - auc\_2: 0.6229 - val\_loss: 0.6695 - val\_accuracy: 0.5939 - val\_auc\_2: 0.6332

Epoch 39/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6681 - accuracy: 0.5841 - auc\_2: 0.6234 - val\_loss: 0.6681 - val\_accuracy: 0.6015 - val\_auc\_2: 0.6337



Epoch 40/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6679 - accuracy: 0.5843 - auc\_2: 0.6236 - val\_loss: 0.6679 - val\_accuracy: 0.5993 - val\_auc\_2: 0.6342

Epoch 41/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6676 - accuracy: 0.5854 - auc\_2: 0.6244 - val\_loss: 0.6694 - val\_accuracy: 0.5944 - val\_auc\_2: 0.6341

Epoch 42/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6676 - accuracy: 0.5841 - auc\_2: 0.6248 - val\_loss: 0.6666 - val\_accuracy: 0.5983 - val\_auc\_2: 0.6349

Epoch 43/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6674 - accuracy: 0.5870 - auc\_2: 0.6253 - val\_loss: 0.6690 - val\_accuracy: 0.5949 - val\_auc\_2: 0.6348

Epoch 44/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6673 - accuracy: 0.5870 - auc\_2: 0.6250 - val\_loss: 0.6679 - val\_accuracy: 0.5985 - val\_auc\_2: 0.6349

Epoch 45/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6669 - accuracy: 0.5868 - auc\_2: 0.6264 - val\_loss: 0.6666 - val\_accuracy: 0.5971 - val\_auc\_2: 0.6354

Epoch 46/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6669 - accuracy: 0.5863 - auc\_2: 0.6259 - val\_loss: 0.6664 - val\_accuracy: 0.5976 - val\_auc\_2: 0.6358

Epoch 47/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6666 - accuracy: 0.5877 - auc\_2: 0.6271 - val\_loss: 0.6682 - val\_accuracy: 0.5959 - val\_auc\_2: 0.6353

Epoch 48/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6670 - accuracy: 0.5870 - auc\_2: 0.6261 - val\_loss: 0.6663 - val\_accuracy: 0.5988 - val\_auc\_2: 0.6359

Epoch 49/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6663 - accuracy: 0.5882 - auc\_2: 0.6280 - val\_loss: 0.6659 - val\_accuracy: 0.6002 - val\_auc\_2: 0.6364

Epoch 50/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6662 - accuracy: 0.5866 - auc\_2: 0.6279 - val\_loss: 0.6663 - val\_accuracy: 0.5983 - val\_auc\_2: 0.6364

Epoch 51/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6658 - accuracy: 0.5867 - auc\_2: 0.6291 - val\_loss: 0.6675 - val\_accuracy: 0.5978 - val\_auc\_2: 0.6363

Epoch 52/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6656 - accuracy: 0.5906 - auc\_2: 0.6294 - val\_loss: 0.6665 - val\_accuracy: 0.5995 - val\_auc\_2: 0.6368

Epoch 53/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6656 - accuracy: 0.5898 - auc\_2: 0.6296 - val\_loss: 0.6669 - val\_accuracy: 0.5966 - val\_auc\_2: 0.6371

Epoch 54/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6654 - accuracy: 0.5893 - auc\_2: 0.6300 - val\_loss: 0.6676 - val\_accuracy: 0.5951 - val\_auc\_2: 0.6370

Epoch 55/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6652 - accuracy: 0.5916 - auc\_2: 0.6311 - val\_loss: 0.6663 - val\_accuracy: 0.6007 - val\_auc\_2: 0.6370

Epoch 56/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6650 - accuracy: 0.5906 - auc\_2: 0.6312 - val\_loss: 0.6655 - val\_accuracy: 0.5995 - val\_auc\_2: 0.6370

Epoch 57/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6647 - accuracy: 0.5920 - auc\_2: 0.6318 - val\_loss: 0.6660 - val\_accuracy: 0.5998 - val\_auc\_2: 0.6371

Epoch 58/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6646 - accuracy: 0.5950 - auc\_2: 0.6319 - val\_loss: 0.6658 - val\_accuracy: 0.6000 - val\_auc\_2: 0.6372

Epoch 59/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6646 - accuracy: 0.5903 - auc\_2: 0.6323 - val\_loss: 0.6663 - val\_accuracy: 0.6000 - val\_auc\_2: 0.6375

Epoch 60/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6644 - accuracy: 0.5953 - auc\_2: 0.6326 - val\_loss: 0.6655 - val\_accuracy: 0.5993 - val\_auc\_2: 0.6375

Epoch 61/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6642 - accuracy: 0.5915 - auc\_2: 0.6329 - val\_loss: 0.6678 - val\_accuracy: 0.5939 - val\_auc\_2: 0.6376

Epoch 62/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6639 - accuracy: 0.5928 - auc\_2: 0.6339 - val\_loss: 0.6647 - val\_accuracy: 0.6007 - val\_auc\_2: 0.6382

Epoch 63/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6641 - accuracy: 0.5922 - auc\_2: 0.6335 - val\_loss: 0.6647 - val\_accuracy: 0.6015 - val\_auc\_2: 0.6381

Epoch 64/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6639 - accuracy: 0.5921 - auc\_2: 0.6338 - val\_loss: 0.6646 - val\_accuracy: 0.6015 - val\_auc\_2: 0.6380

Epoch 65/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6637 - accuracy: 0.5944 - auc\_2: 0.6344 - val\_loss: 0.6657 - val\_accuracy: 0.6002 - val\_auc\_2: 0.6382

Epoch 66/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6638 - accuracy: 0.5929 - auc\_2: 0.6339 - val\_loss: 0.6670 - val\_accuracy: 0.5959 - val\_auc\_2: 0.6381

Epoch 67/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6635 - accuracy: 0.5918 - auc\_2: 0.6344 - val\_loss: 0.6647 - val\_accuracy: 0.6024 - val\_auc\_2: 0.6384

Epoch 68/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6632 - accuracy: 0.5941 - auc\_2: 0.6361 - val\_loss: 0.6643 - val\_accuracy: 0.6012 - val\_auc\_2: 0.6384

Epoch 69/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6632 - accuracy: 0.5948 - auc\_2: 0.6360 - val\_loss: 0.6654 - val\_accuracy: 0.5998 - val\_auc\_2: 0.6383

Epoch 70/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6630 - accuracy: 0.5950 - auc\_2: 0.6359 - val\_loss: 0.6648 - val\_accuracy: 0.6002 - val\_auc\_2: 0.6385

Epoch 71/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6629 - accuracy: 0.5942 - auc\_2: 0.6363 - val\_loss: 0.6656 - val\_accuracy: 0.6012 - val\_auc\_2: 0.6385

Epoch 72/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6628 - accuracy: 0.5951 - auc\_2: 0.6364 - val\_loss: 0.6656 - val\_accuracy: 0.5998 - val\_auc\_2: 0.6385

Epoch 73/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6626 - accuracy: 0.5943 - auc\_2: 0.6369 - val\_loss: 0.6680 - val\_accuracy: 0.5932 - val\_auc\_2: 0.6386

Epoch 74/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6626 - accuracy: 0.5945 - auc\_2: 0.6367 - val\_loss: 0.6655 - val\_accuracy: 0.6002 - val\_auc\_2: 0.6387

Epoch 75/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6625 - accuracy: 0.5957 - auc\_2: 0.6374 - val\_loss: 0.6654 - val\_accuracy: 0.6015 - val\_auc\_2: 0.6390

Epoch 76/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6624 - accuracy: 0.5950 - auc\_2: 0.6375 - val\_loss: 0.6640 - val\_accuracy: 0.6012 - val\_auc\_2: 0.6393

Epoch 77/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6625 - accuracy: 0.5959 - auc\_2: 0.6373 - val\_loss: 0.6654 - val\_accuracy: 0.6015 - val\_auc\_2: 0.6390

Epoch 78/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6623 - accuracy: 0.5930 - auc\_2: 0.6380 - val\_loss: 0.6646 - val\_accuracy: 0.6019 - val\_auc\_2: 0.6396

Epoch 79/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6621 - accuracy: 0.5952 - auc\_2: 0.6376 - val\_loss: 0.6642 - val\_accuracy: 0.6007 - val\_auc\_2: 0.6394

Epoch 80/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6618 - accuracy: 0.5963 - auc\_2: 0.6383 - val\_loss: 0.6657 - val\_accuracy: 0.5988 - val\_auc\_2: 0.6388

Epoch 81/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6618 - accuracy: 0.5942 - auc\_2: 0.6387 - val\_loss: 0.6642 - val\_accuracy: 0.6005 - val\_auc\_2: 0.6391

Epoch 82/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6622 - accuracy: 0.5972 - auc\_2: 0.6376 - val\_loss: 0.6657 - val\_accuracy: 0.5988 - val\_auc\_2: 0.6390

Epoch 83/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6617 - accuracy: 0.5944 - auc\_2: 0.6393 - val\_loss: 0.6638 - val\_accuracy: 0.6019 - val\_auc\_2: 0.6397

Epoch 84/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6615 - accuracy: 0.5970 - auc\_2: 0.6392 - val\_loss: 0.6640 - val\_accuracy: 0.5995 - val\_auc\_2: 0.6397

Epoch 85/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6615 - accuracy: 0.5964 - auc\_2: 0.6395 - val\_loss: 0.6646 - val\_accuracy: 0.6027 - val\_auc\_2: 0.6394

Epoch 86/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6614 - accuracy: 0.5983 - auc\_2: 0.6399 - val\_loss: 0.6652 - val\_accuracy: 0.5998 - val\_auc\_2: 0.6393

Epoch 87/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6613 - accuracy: 0.5963 - auc\_2: 0.6398 - val\_loss: 0.6638 - val\_accuracy: 0.6005 - val\_auc\_2: 0.6397

Epoch 88/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6613 - accuracy: 0.5950 - auc\_2: 0.6401 - val\_loss: 0.6643 - val\_accuracy: 0.6010 - val\_auc\_2: 0.6396

Epoch 89/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6611 - accuracy: 0.5975 - auc\_2: 0.6403 - val\_loss: 0.6639 - val\_accuracy: 0.5995 - val\_auc\_2: 0.6397

Epoch 90/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6610 - accuracy: 0.5981 - auc\_2: 0.6409 - val\_loss: 0.6645 - val\_accuracy: 0.6012 - val\_auc\_2: 0.6394

Epoch 91/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6608 - accuracy: 0.5963 - auc\_2: 0.6410 - val\_loss: 0.6638 - val\_accuracy: 0.5995 - val\_auc\_2: 0.6399

Epoch 92/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6607 - accuracy: 0.5981 - auc\_2: 0.6415 - val\_loss: 0.6650 - val\_accuracy: 0.6010 - val\_auc\_2: 0.6398

Epoch 93/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6609 - accuracy: 0.5983 - auc\_2: 0.6405 - val\_loss: 0.6646 - val\_accuracy: 0.6015 - val\_auc\_2: 0.6400

Epoch 94/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6608 - accuracy: 0.5979 - auc\_2: 0.6411 - val\_loss: 0.6641 - val\_accuracy: 0.6010 - val\_auc\_2: 0.6398

Epoch 95/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6605 - accuracy: 0.5976 - auc\_2: 0.6417 - val\_loss: 0.6641 - val\_accuracy: 0.6019 - val\_auc\_2: 0.6400

Epoch 96/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6605 - accuracy: 0.5976 - auc\_2: 0.6416 - val\_loss: 0.6630 - val\_accuracy: 0.6019 - val\_auc\_2: 0.6401

Epoch 97/100  
96/96 [=====] - 0s 4ms/step - loss: 0.6603 - accuracy: 0.5989 - auc\_2: 0.6420 - val\_loss: 0.6630 - val\_accuracy: 0.6034 - val\_auc\_2: 0.6400

Epoch 98/100  
96/96 [=====] - 1s 5ms/step - loss: 0.6602 - accuracy: 0.5972 - auc\_2: 0.6425 - val\_loss: 0.6633 - val\_accuracy: 0.6029 - val\_auc\_2: 0.6400

Epoch 99/100  
96/96 [=====] - 0s 5ms/step - loss: 0.6603 - accuracy: 0.5972 - auc\_2: 0.6421 - val\_loss: 0.6648 - val\_accuracy: 0.5998 - val\_auc\_2: 0.6398


Epoch 100/100

96/96 [=====] - 0s 5ms/step - loss: 0.6603 - accuracy: 0.5984 - auc\_2: 0.6426 - val\_loss: 0.6642 - val\_accuracy: 0.6022 - val\_auc\_2: 0.6401

### Q3.7 assessing model performance

Make 3 plots: one each for loss, accuracy and AUC. Each plot should have train and validation scores labeled.

```
[20]: plt.plot(hist.history['loss'])
plt.plot(hist.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```




```
[21]: # YOUR CODE HERE #
print(hist.history)

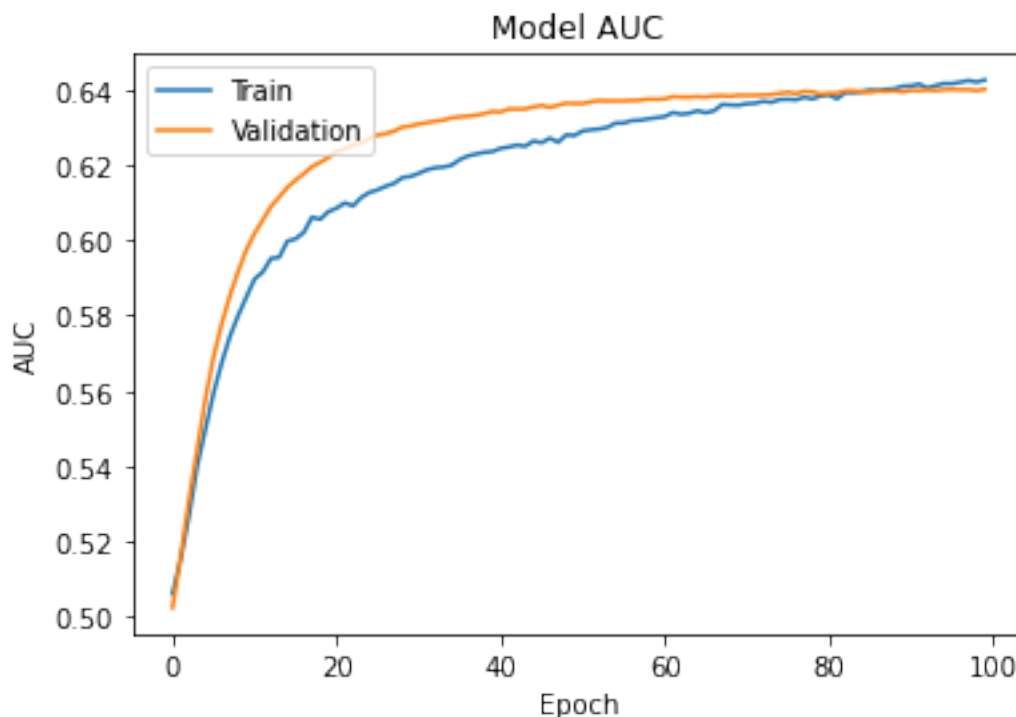
plt.plot(hist.history['accuracy'])
plt.plot(hist.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
```

```
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# END CODE #
```



```
[22]: plt.plot(hist.history['auc_2'])
plt.plot(hist.history['val_auc_2'])
plt.title('Model AUC')
plt.ylabel('AUC')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



### Q3.8 combining chunked predictions to full predictions for readmission

In the [ClinicalBert paper](#), the authors did pretraining with MIMIC-III, and in section 3.3.2, they make predictions of hospital readmission on MIMIC-III, which is the same task we are doing.

We've broken our notes into chunks and made predictions for each chunk. The ClinicalBert authors propose a method to combine the chunked notes predictions into one prediction per note (equation 4 in the paper).

Implement this function in the next cell for the *validation* set with `c=1`. Use the combined predictions to compute the AUC.

```
[23]: import sklearn
def predict_FULL_note_readmission_clinicalBert(model,
bert_pool_embeddings_CHUNKS,
idxs_CHUNKS, c=1):
    """
    Combine CHUNK predictions to FULL predictions using equation 4 of
    https://arxiv.org/pdf/1904.05342.pdf

    Args:
    model (tf.keras.Model): a trained prediction model.
    bert_embedding_CHUNK (np.array[float,flat]): chunked dataset of bert
    embeddings
    for running prediction.
```



```

    idxs_CHUNKS np.array([int]): idxs_CHUNKS[i]=j means chunk i is a
    ↳subsequence of
        a note i.

Returns:
    y_pred_FULL (np.array([int]))
    """
    n_unique = len(np.unique(idxs_CHUNKS))
    idxs_CHUNKS = np.array(idxs_CHUNKS)
    y_pred_score_CHUNKS = model.predict(bert_pool_embeddings_CHUNKS)

    y_pred_FULL = np.zeros(n_unique)
    for i in range(n_unique):
        # YOUR CODE HERE #
        notes = np.where(idxs_CHUNKS == i)
        n = len(notes)
        chunks = y_pred_score_CHUNKS[notes]
        max_chunk = max(chunks)
        avg_chunks = np.average(chunks)
        readmit = (avg_chunks + (max_chunk[0] * (n/c))/(1 + (n/c)))
        y_pred_FULL[i] = readmit
        # END CODE #
    #y_pred_FULL = y_pred_FULL[y_pred_FULL != 0]
    #y_pred_FULL = sklearn.preprocessing.binarize(y_pred_FULL.reshape(-1, 1), 0.
    ↳5)
    return y_pred_FULL

c = 1
y_pred_FULL = predict_FULL_note_readmission_clinicalBert(model,

    ↳val_bert_pool_embeddings_CHUNKS,
                                                val_idxes_CHUNKS,
                                                c=c)

y_score_FULL = val_labels_FULL
auc = None
# YOUR CODE HERE #
#auc = sklearn.metrics.roc_auc_score(y_pred_FULL, y_score_FULL)
auc = sklearn.metrics.roc_auc_score(y_score_FULL, y_pred_FULL)
#sklearn.metrics.roc_auc_score(y_true, y_score
# END CODE #
print(f"AUC {auc:.5f}")

```

AUC 0.69830

### Q3.9 hyperparameter tuning

Run `predict_FULL_note_readmission_clinicalBert` and compute the AUC for a range of `c` values. We will use the best AUC to choose a value of `c`. This is hyperparameter tuning, and so

we should do this on the validation set.

```
[24]: for c in [0.01,0.1,0.5,1,2,5,10,20,50]:
    auc = None
    # YOUR CODE HERE #
    y_pred_FULL = predict_FULL_note_readmission_clinicalBert(model,
                                                                val_idxes_CHUNKS,
                                                                c=c)
    auc = sklearn.metrics.roc_auc_score(y_score_FULL, y_pred_FULL)
    # END CODE #
    print(f"c {c}\t AUC: {auc:.5f}")
```

```
c 0.01    AUC: 0.69676
c 0.1     AUC: 0.69693
c 0.5     AUC: 0.69799
c 1       AUC: 0.69830
c 2       AUC: 0.69836
c 5       AUC: 0.69789
c 10      AUC: 0.69756
c 20      AUC: 0.69694
c 50      AUC: 0.69648
```

### Q3.10 evaluation

Now that you have chosen a *c* value, let's evaluate on the test set. Run `predict_FULL_note_readmission_clinicalBert` and compute the AUC. In the next cell you just have to fill in your value of *c* and compute the auc.

```
[25]: c = 2 # choose your best value

y_pred_FULL = predict_FULL_note_readmission_clinicalBert(model,
                                                            test_idxes_CHUNKS,
                                                            c=c)

y_score_FULL = test_labels_FULL
auc = None
# YOUR CODE HERE #
auc = sklearn.metrics.roc_auc_score(y_score_FULL, y_pred_FULL)
# END CODE #
print(f"Test set auc {auc:.5f}")
```

Test set auc 0.68074

Your test set AUC may be different to the validation set. Explain why, and give one strategy for getting more consistent results between validation and test.

*Written answer:* randomly generated test set was easier to predict than the validation as the model

was pre-trained on words that were more relevant to the test set. try and normalize btwn the sets.

we should do this on the validation set.

```
[24]: for c in [0.01,0.1,0.5,1,2,5,10,20,50]:
    auc = None
    # YOUR CODE HERE #
    y_pred_FULL = predict_FULL_note_readmission_clinicalBert(model,
                                                                val_bert_pool_embeddings_CHUNKS,
                                                                val_idxes_CHUNKS,
                                                                c=c)
    auc = sklearn.metrics.roc_auc_score(y_score_FULL, y_pred_FULL)
    # END CODE #
    print(f"c {c}\t AUC: {auc:.5f}")
```

```
c 0.01    AUC: 0.69676
c 0.1     AUC: 0.69693
c 0.5     AUC: 0.69799
c 1       AUC: 0.69830
c 2       AUC: 0.69836
c 5       AUC: 0.69789
c 10      AUC: 0.69756
c 20      AUC: 0.69694
c 50      AUC: 0.69648
```

### Q3.10 evaluation

Now that you have chosen a *c* value, let's evaluate on the test set. Run `predict_FULL_note_readmission_clinicalBert` and compute the AUC. In the next cell you just have to fill in your value of *c* and compute the auc.

```
[25]: c = 2 # choose your best value

y_pred_FULL = predict_FULL_note_readmission_clinicalBert(model,
                                                            test_bert_pool_embeddings_CHUNKS,
                                                            test_idxes_CHUNKS,
                                                            c=c)

y_score_FULL = test_labels_FULL
auc = None
# YOUR CODE HERE #
auc = sklearn.metrics.roc_auc_score(y_score_FULL, y_pred_FULL)
# END CODE #
print(f"Test set auc {auc:.5f}")
```

Test set auc 0.68074

Your test set AUC may be different to the validation set. Explain why, and give one strategy for getting more consistent results between validation and test.

*Written answer:* The randomly generated test set was more difficult to predict than the validation

as the model was pre-trained on words that were more relevant to the validation set. A suitable fix to this issue to get more consistent results between validation and test is to try and normalize between both datasets and ensure that words in the validation and test set notes chunks had relatively the same frequencies. As a result, we would see less imbalance between the test and validation set and see more consistent results.

[ ]: